



---

## ***Chap. 8) Memory Management Strategies***

경희대학교 컴퓨터공학과

방 재 훈

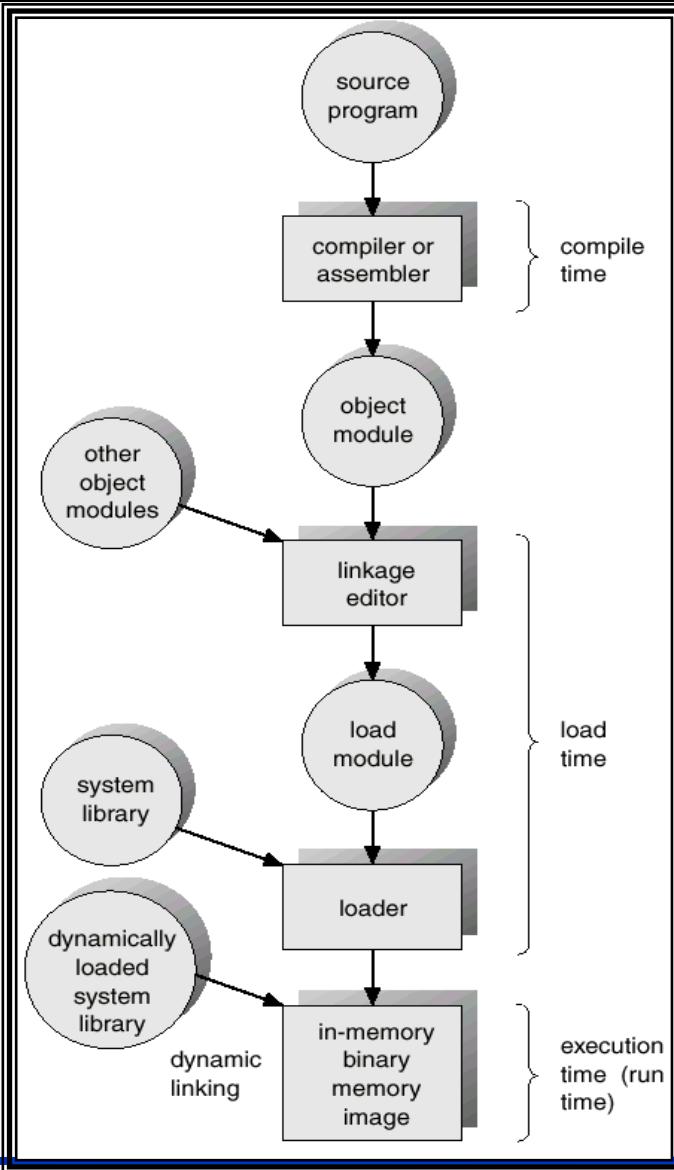
# *Binding of Instructions and Data to Memory*

---

- Address binding of instructions and data to memory addresses can happen at three different stages
  - ✓ Compile time
    - If memory location known a priori, absolute code can be generated
    - must recompile code if starting location changes
  - ✓ Load time
    - Must generate *relocatable* code if memory location is not known at compile time
  - ✓ Execution time
    - Binding delayed until run time if the process can be moved during its execution from one memory segment to another
    - Need hardware support for address maps (e.g., *base* and *limit registers*)



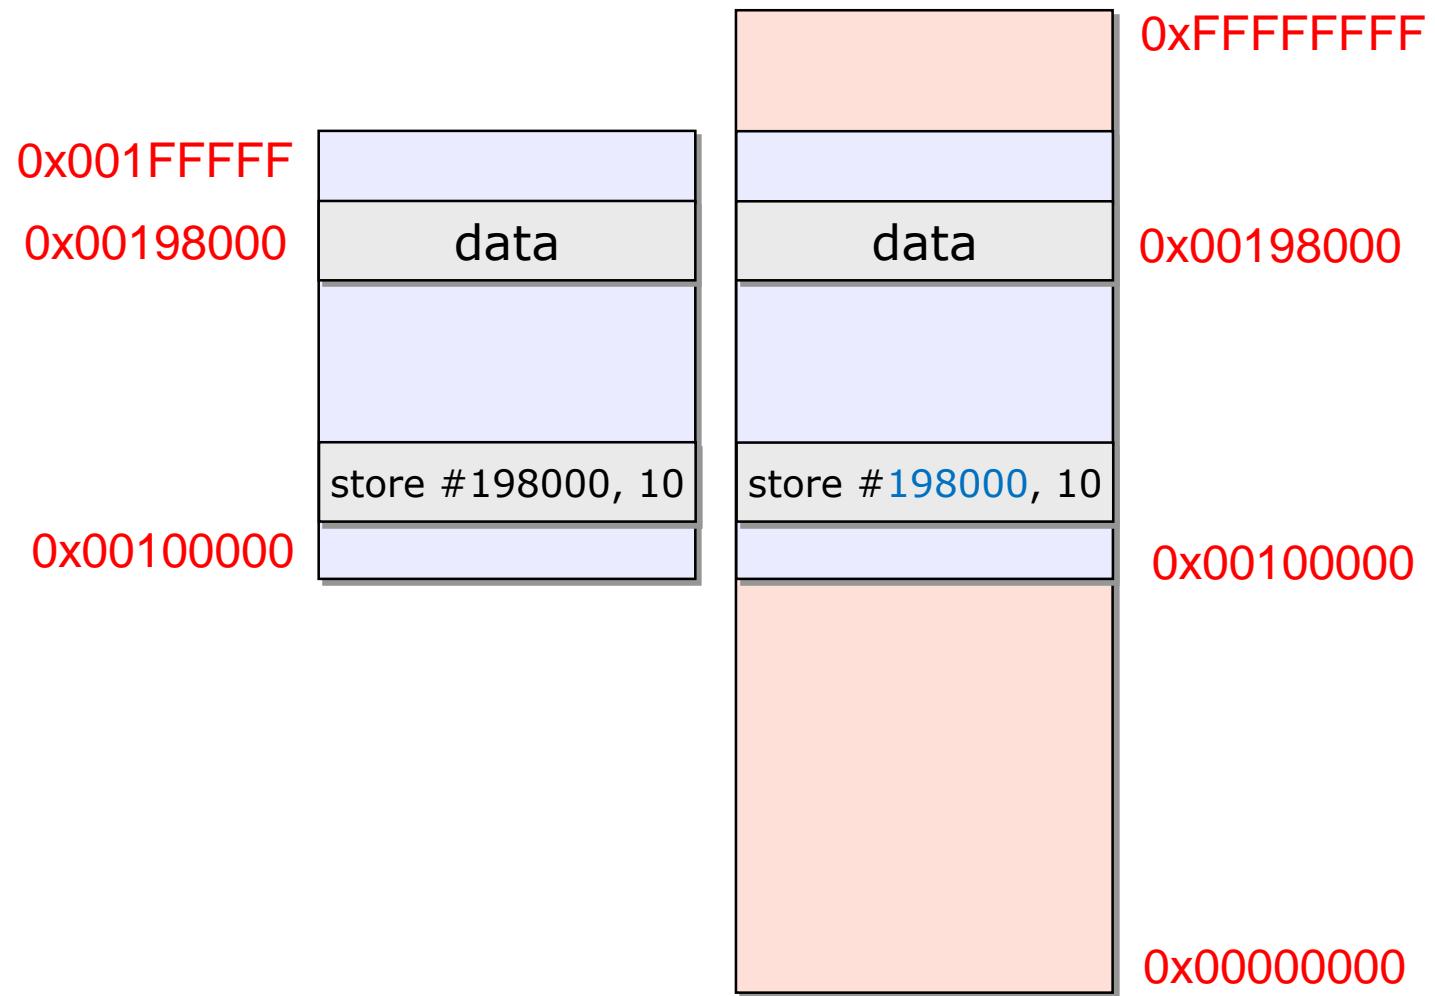
# Multistep Processing of a User Program



# *Binding of Memory Address*

## ■ Compile time

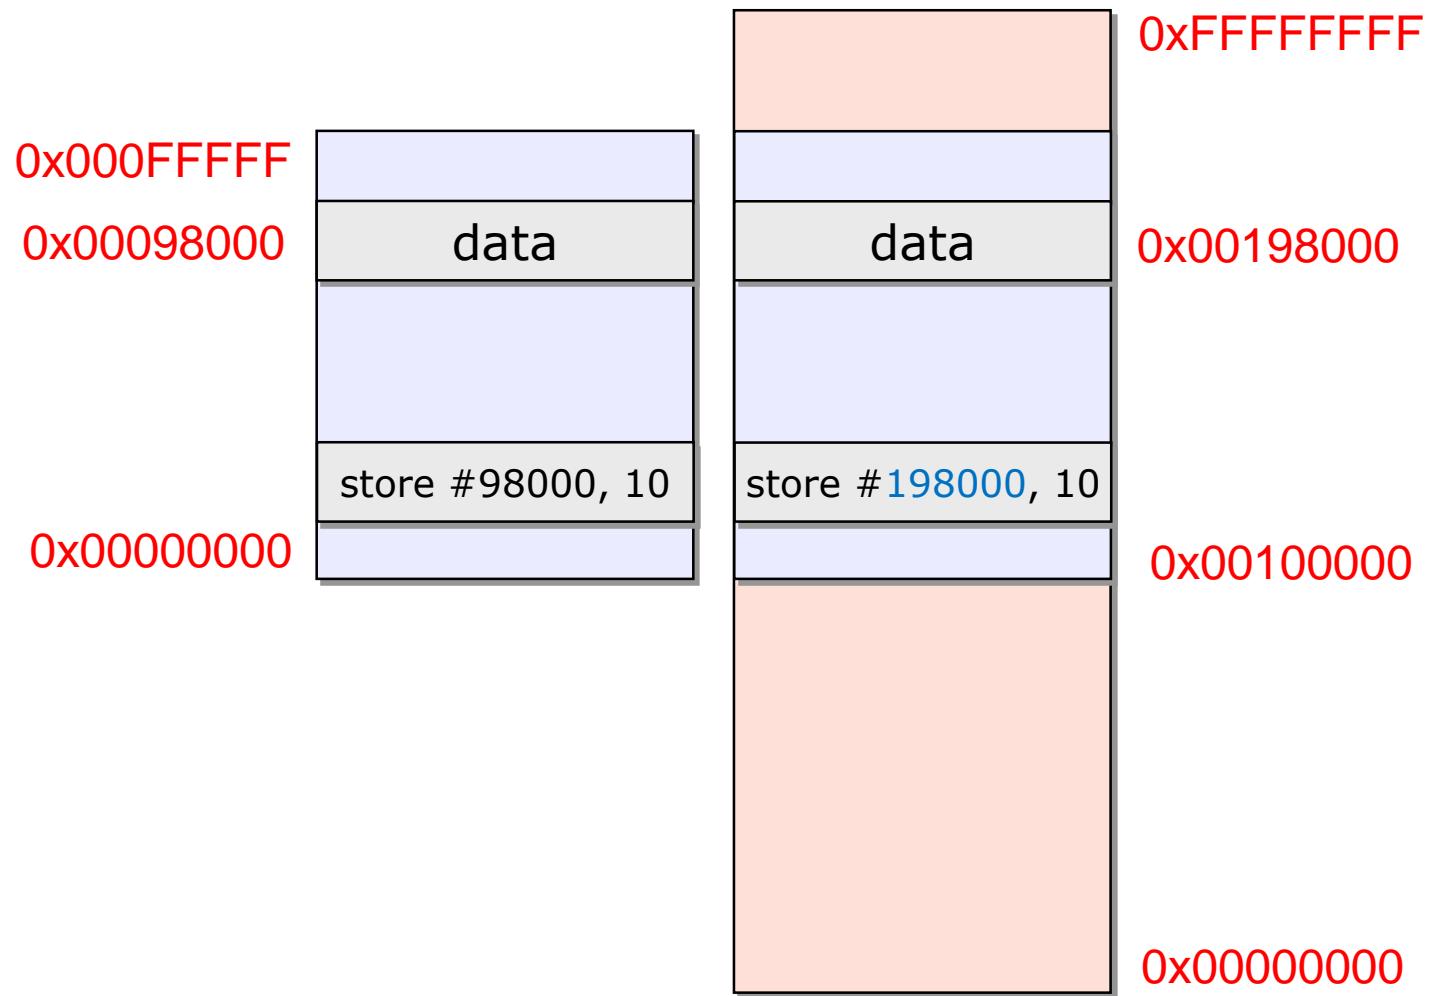
```
void func( )  
{  
    int data;  
    ...  
    data = 10;  
    ...  
}
```



# *Binding of Memory Address*

## ■ Load time

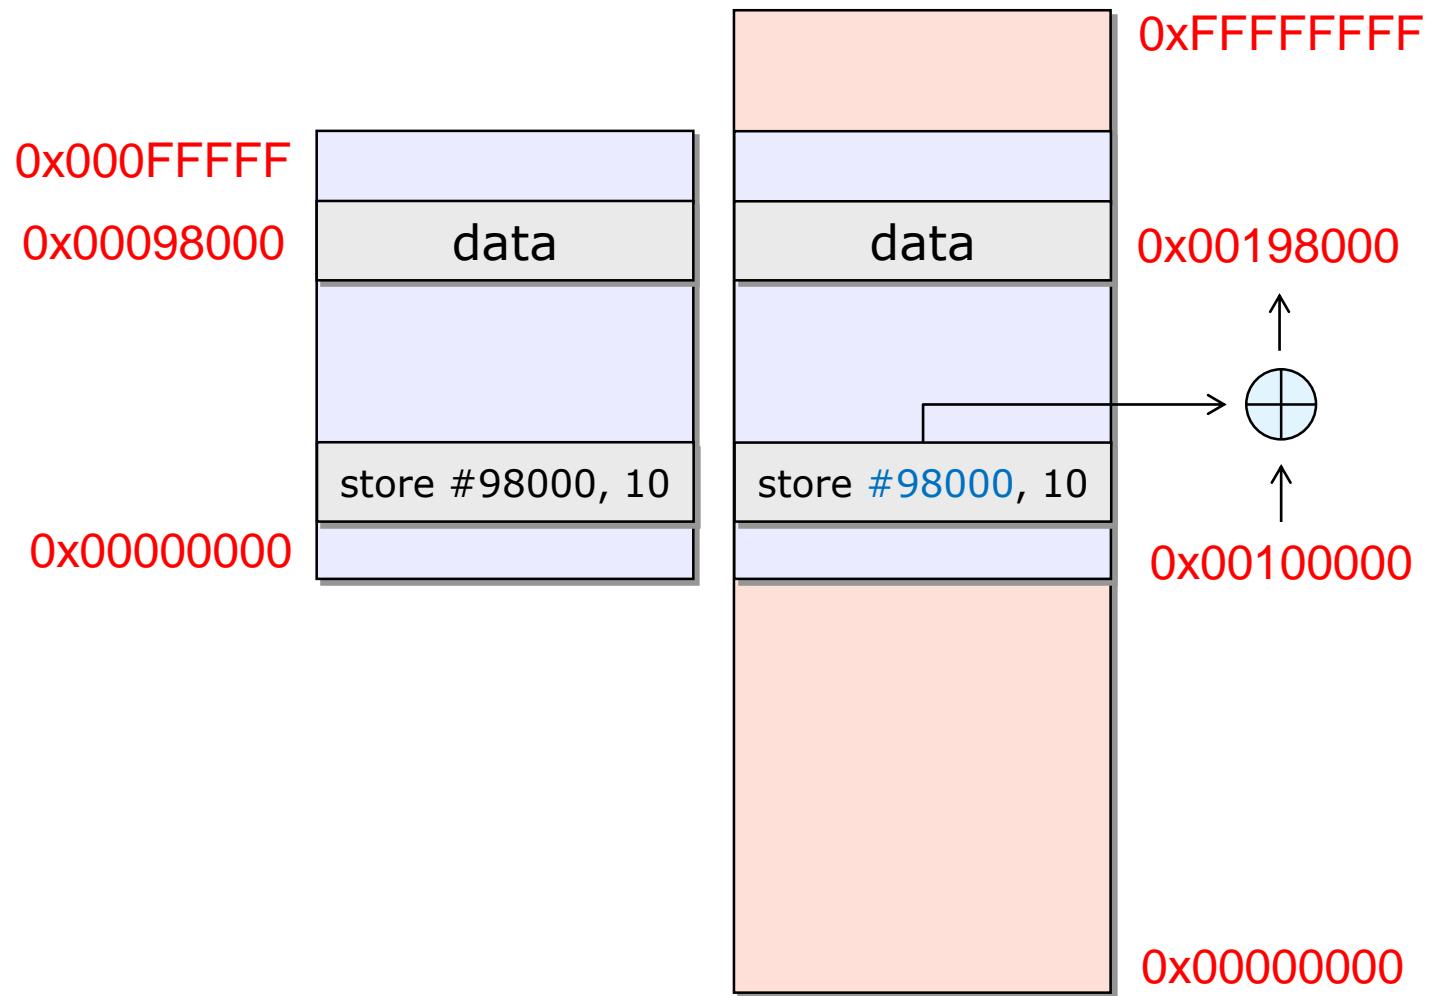
```
void func( )  
{  
    int data;  
    ...  
    data = 10;  
    ...  
}
```



# *Binding of Memory Address*

## ■ Execution time

```
void func( )  
{  
    int data;  
    ...  
    data = 10;  
    ...  
}
```



# *Memory Management*

---

## ■ Goals

- ✓ To provide a convenient abstraction for programming
- ✓ To allocate scarce memory resources among competing processes to maximize performance with minimal overhead
- ✓ To provide isolation between processes



# ***Memory Management (Cont'd)***

---

## ■ Batch programming

- ✓ Programs use physical addresses directly
- ✓ OS loads job, runs it, unloads it

## ■ Multiprogramming

- ✓ Need multiple processes in memory at once
  - Overlap I/O and CPU of multiple jobs
- ✓ Can do it a number of ways
  - Fixed and variable partitioning, paging, segmentation, etc.
- ✓ Requirements
  - Protection: restrict which addresses processes can use
  - Fast translation: memory lookups must be fast, in spite of protection scheme
  - Fast context switching: updating memory hardware (for protection and translation) should be quick



# ***Memory Management (Cont'd)***

---

## ■ Issues

- ✓ Support for multiple processes
  - Each process should have a logically contiguous space
  - The size of each space is variable
- ✓ Enable a process to be larger than the amount of memory allocated to it
  - Not all the memory spaces are used simultaneously
  - Memory references exhibit spatial and temporal locality
- ✓ Protection
- ✓ Sharing
- ✓ Support for multiple regions per process (segments)
- ✓ Performance
  - Memory reference overhead
  - Context switching overhead



# **Memory Management (Cont'd)**

---

## ■ Solution: Virtual Memory (VM)

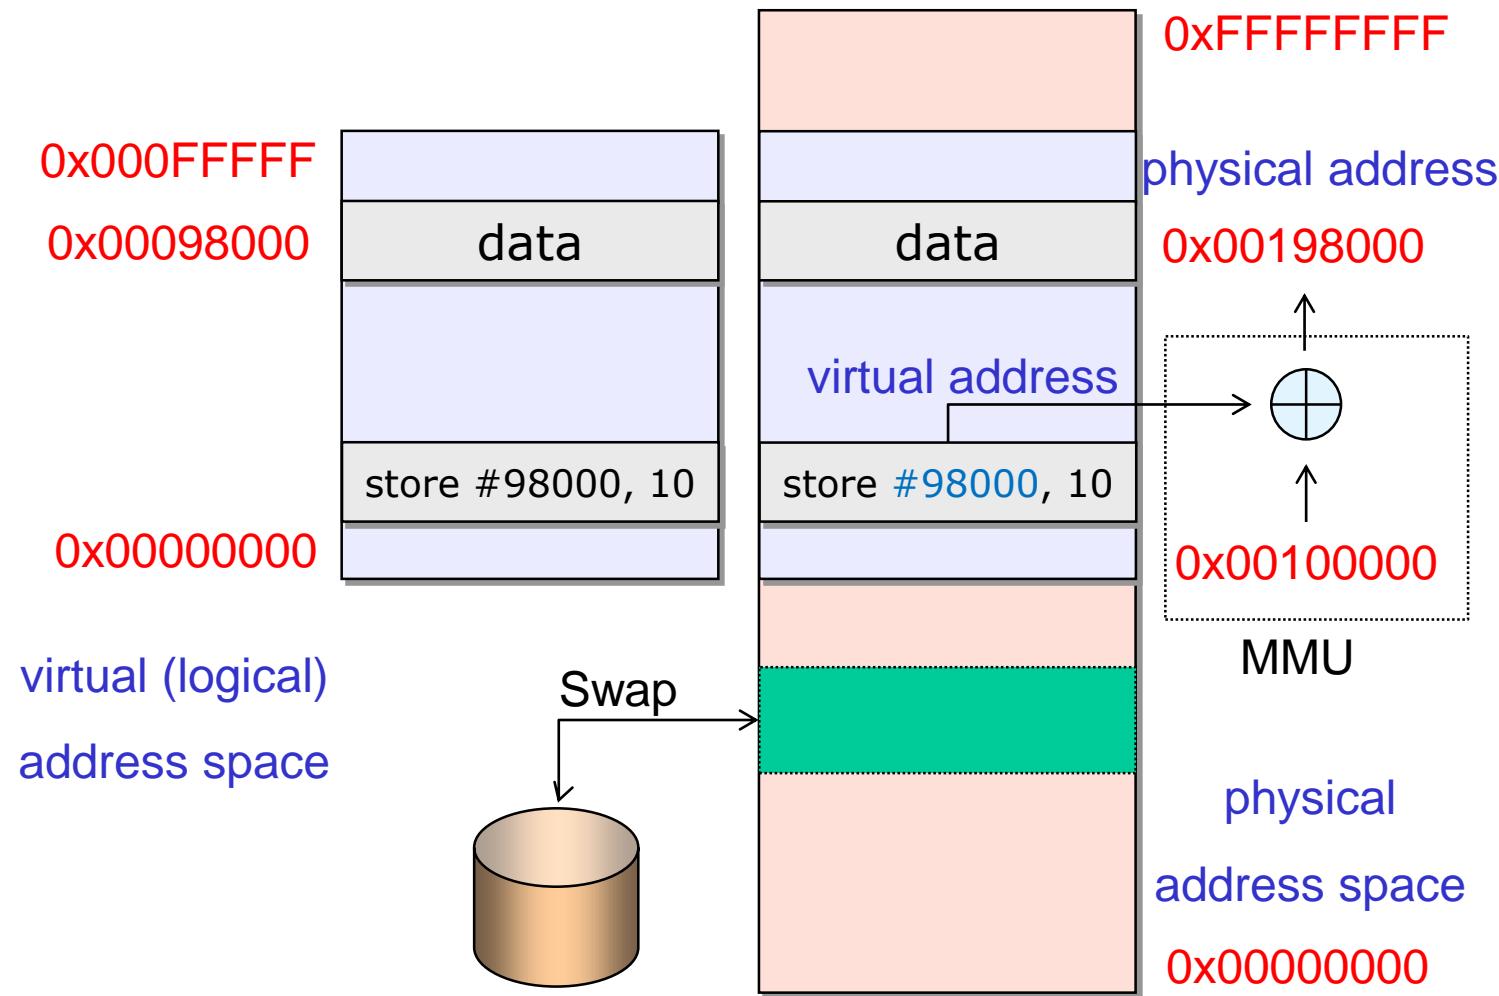
- ✓ VM enables programs to execute without requiring their entire address space to be resident in physical memory
  - Program can also execute on machines with less RAM than it needs
- ✓ Many programs don't need all of their code or data at once (or ever)
  - e.g., branches they never take, or data they never R/W
  - No need to allocate memory for it, OS should adjust amount allocated based on its run-time behavior
- ✓ VM isolates processes from each other
  - One process cannot name addresses visible to others
  - Each process has its own isolated address space



# Memory Management (Cont'd)

## Virtual memory management

```
void func( )  
{  
    int data;  
    ...  
    data = 10;  
    ...  
}
```



# *Dynamic Loading*

---

- Routine is not loaded until it is called
- Better memory-space utilization
  - ✓ Unused routine is never loaded
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
  - ✓ Implemented through program design



# ***Dynamic Linking***

---

- Linking postponed until execution time
- Small piece of code, *stub*, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system needed to check if routine is in processes' memory address
- Dynamic linking is particularly useful for libraries



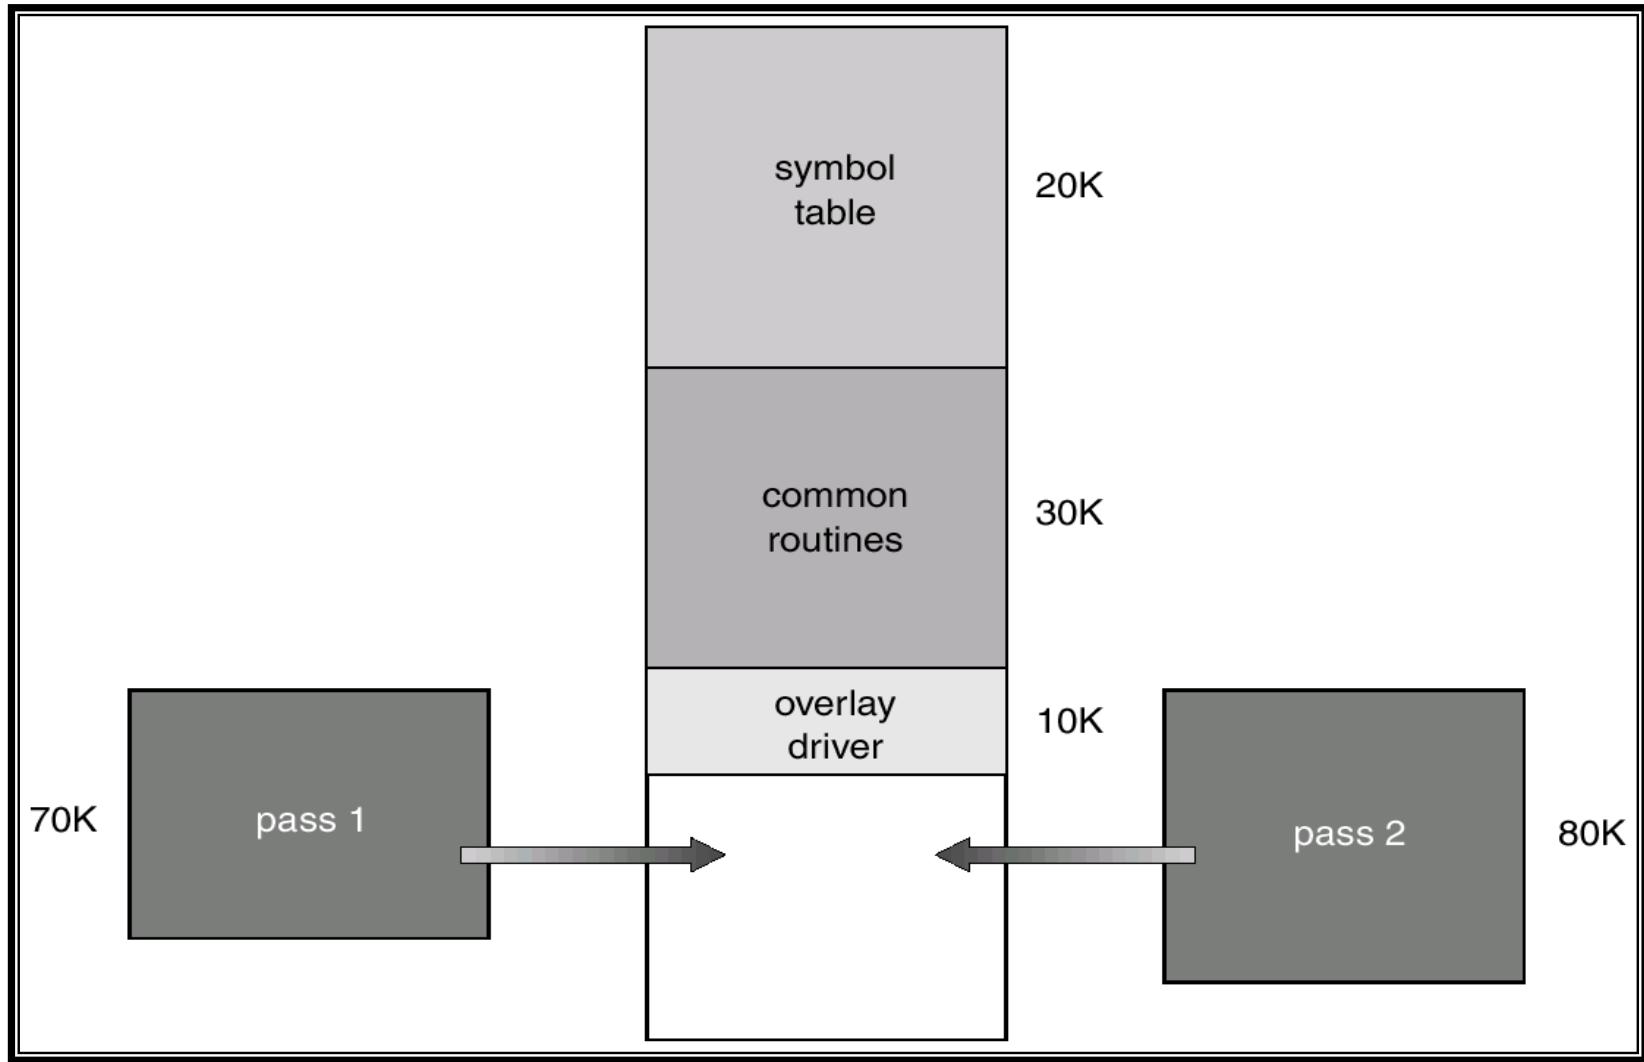
# Overlays

---

- Keep in memory only those instructions and data that are needed at any given time
- Needed when process is larger than amount of memory allocated to it
- Implemented by user, no special support needed from operating system, programming design of overlay structure is complex



# Overlays for a Two-Pass Assembler



# Swapping

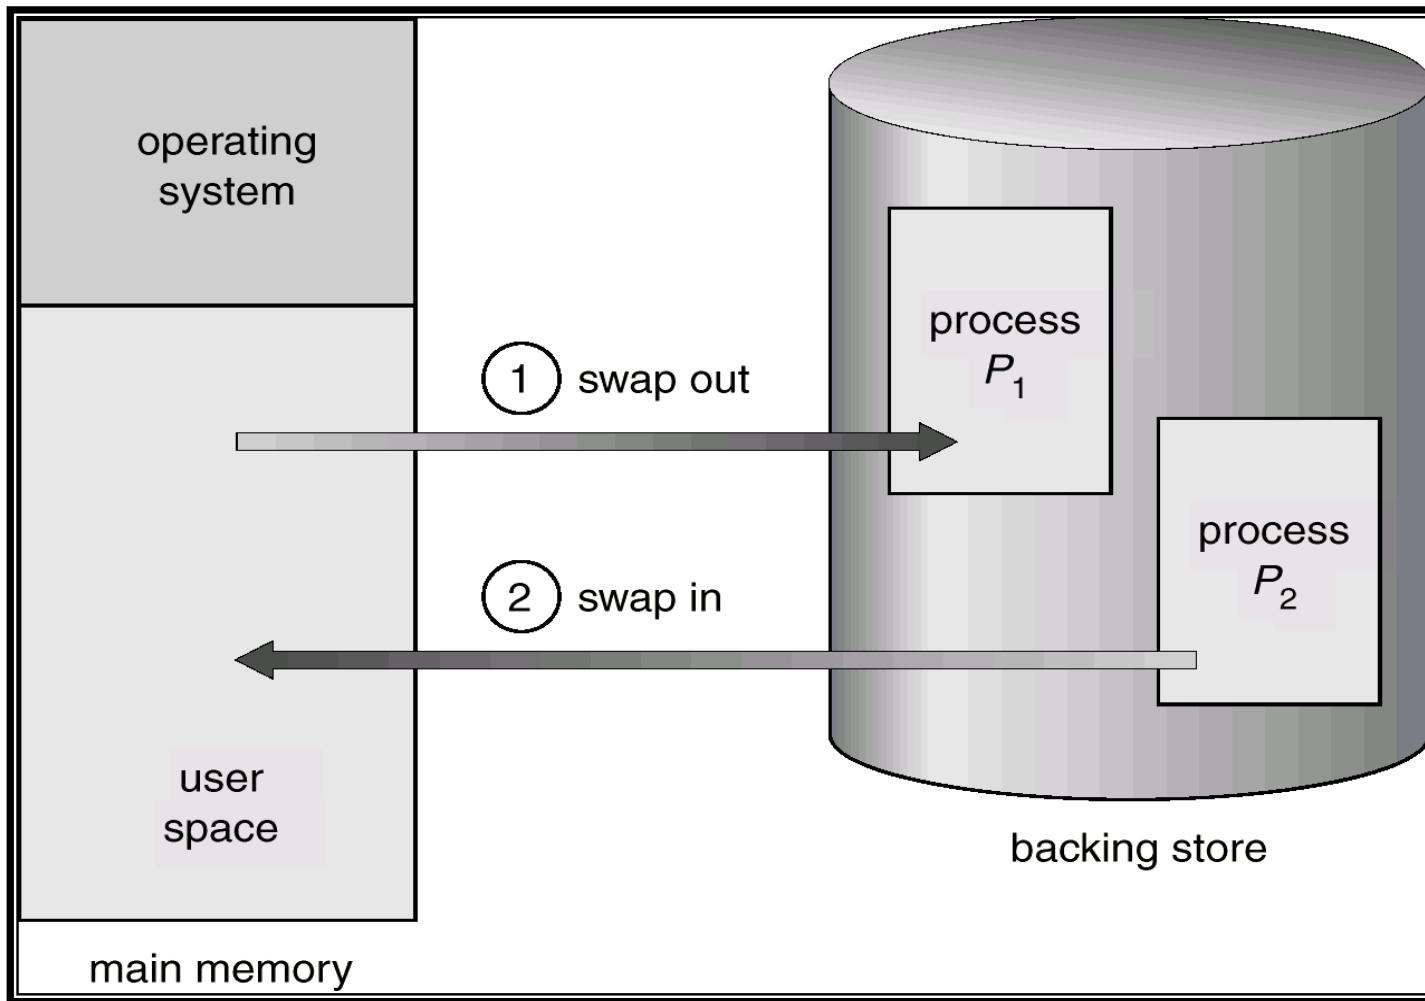
---

- A process can be *swapped* temporarily out of memory to a *Backing store*, and then brought back into memory for continued execution
- Backing store
  - ✓ Fast disk large enough to accommodate copies of all memory images for all users
  - ✓ Must provide direct access to these memory images
- *Roll out & Roll in*
  - ✓ Swapping variant used for priority-based scheduling algorithms
  - ✓ Lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time
  - ✓ Total transfer time is directly proportional to the *amount* of memory swapped
- Modified versions of swapping are found on many systems
  - ✓ i.e., UNIX, Linux, and Windows



# Schematic View of Swapping

---



# *Logical vs. Physical Address Space*

---

- The concept of a logical *address space* that is bound to a separate *physical address space* is central to proper memory management
  - ✓ *Logical address*
    - generated by the CPU
    - also referred to as *virtual address*
  - ✓ *Physical address*
    - address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes
  - ✓ Logical (virtual) and physical addresses differ in execution-time address-binding scheme



# *Logical (Virtual) Address Space*

## ■ Example

```
#include <stdio.h>

int n = 0;

int main ()
{
    printf ("%n = 0x%08x\n", &n);
}

% ./a.out
&n = 0x08049508
% ./a.out
&n = 0x08049508
```

- ✓ What happens if two users simultaneously run this application?



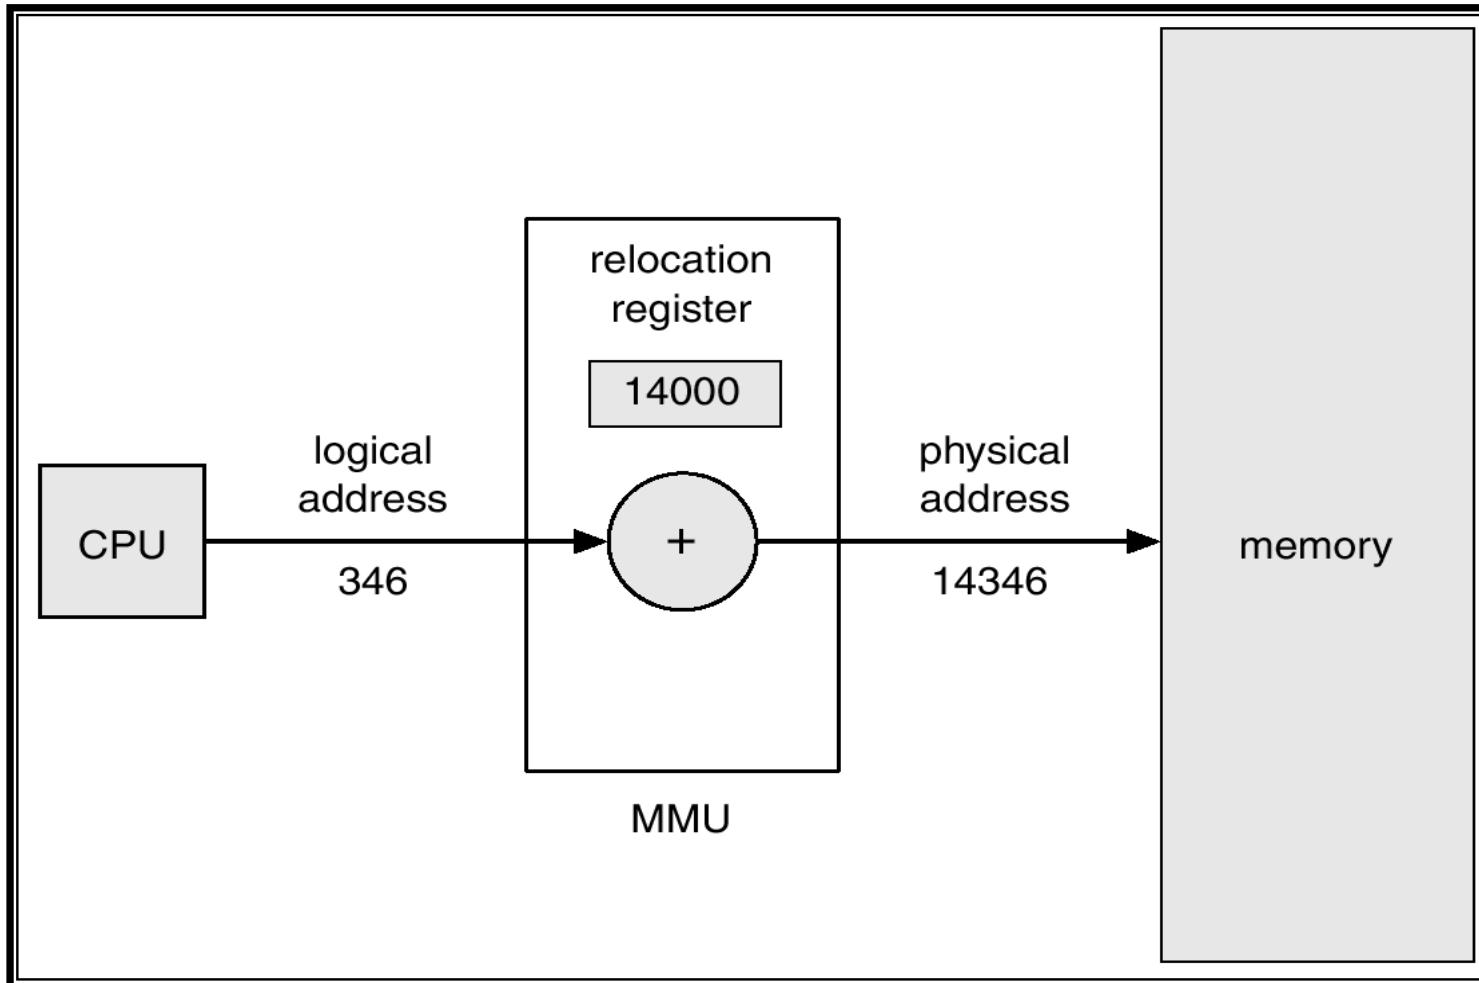
# **Memory-Management Unit (MMU)**

---

- Hardware device that maps logical (virtual) to physical address
- In MMU scheme, the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- The user program deals with *logical* addresses
  - ✓ It never sees the *real* physical addresses



# *Dynamic relocation using a relocation register*



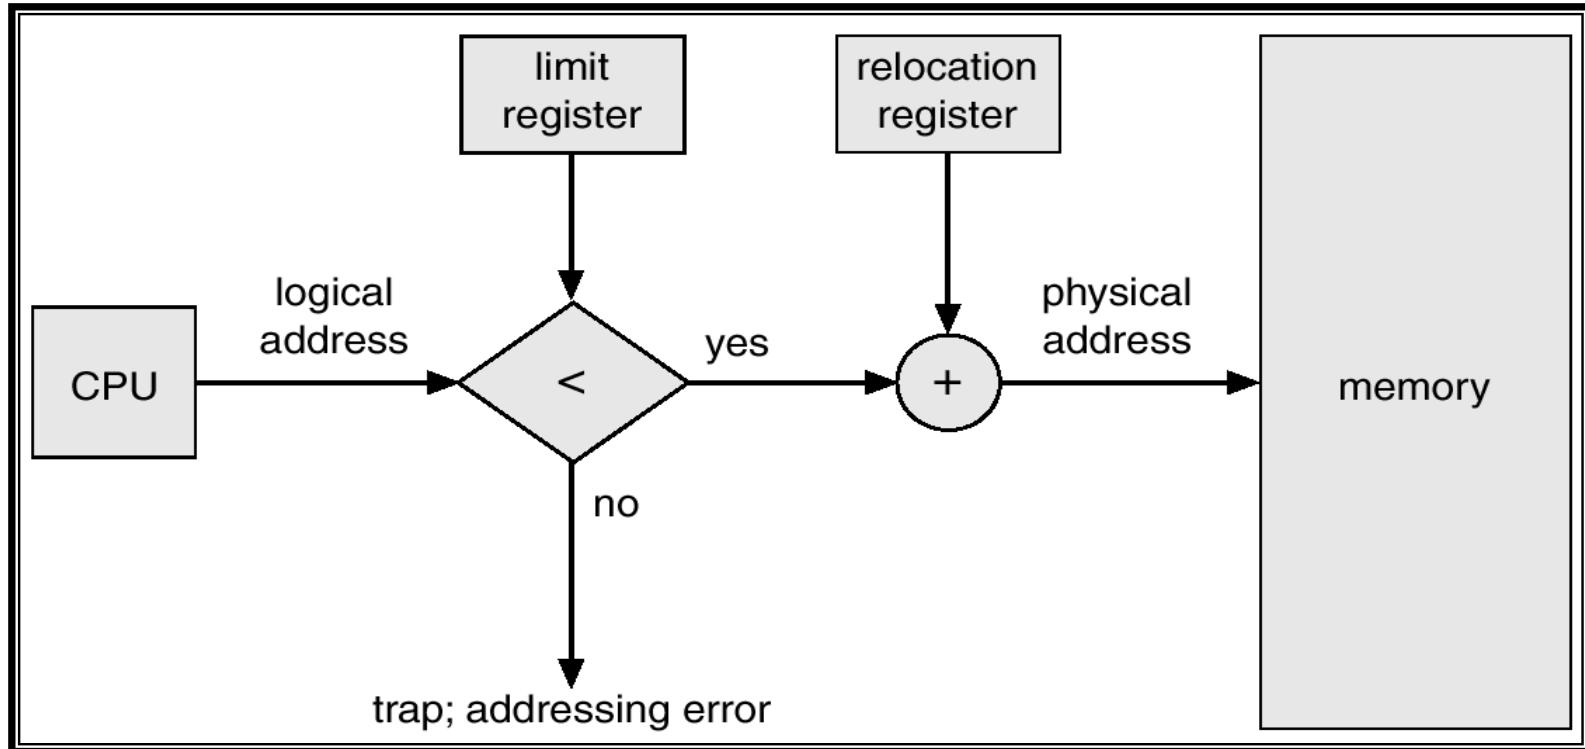
# **Contiguous Allocation**

---

- Main memory usually into two partitions:
  - ✓ Resident operating system, usually held in low memory with interrupt vector
  - ✓ User processes then held in high memory
- Single-partition allocation
  - ✓ Relocation-register scheme used to protect user processes from each other, and from changing operating-system code and data
  - ✓ Relocation register contains value of smallest physical address
  - ✓ Limit register contains range of logical addresses
  - ✓ Each logical address must be less than the limit register
- Cf) *Physically contiguous vs. Logically contiguous allocation*



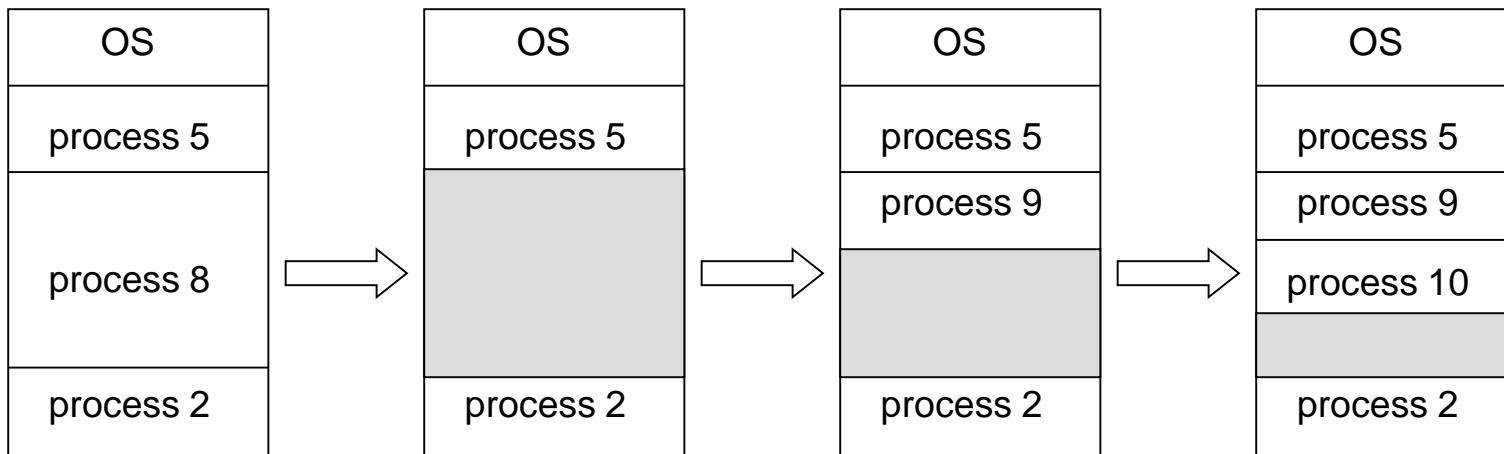
# *Hardware Support for Relocation and Limit Registers*



# Contiguous Allocation (Cont'd)

## Multiple-partition allocation

- ✓ *Hole*
  - block of available memory
  - holes of various size are scattered throughout memory
- ✓ When a process arrives, it is allocated memory from a hole large enough to accommodate it
- ✓ Operating system maintains information about:
  - a) allocated partitions   b) free partitions (hole)



# **Dynamic Storage-Allocation Problem**

---

- How to satisfy a request of size  $n$  from a list of free holes
  - ✓ First-fit
    - Allocate the *first* hole that is big enough
  - ✓ Best-fit
    - Allocate the *smallest* hole that is big enough
    - must search entire list, unless ordered by size
    - Produces the smallest leftover hole
  - ✓ Worst-fit
    - Allocate the *largest* hole
    - must also search entire list
    - Produces the largest leftover hole
- First-fit and best-fit better than worst-fit in terms of speed and storage utilization



# *Fragmentation*

---

## ■ External Fragmentation

- ✓ Total memory space exists to satisfy a request, but it is not contiguous

## ■ Internal Fragmentation

- ✓ Allocated memory may be slightly larger than requested memory
- ✓ This size difference is memory internal to a partition, but not being used

## ■ Reduce external fragmentation by compaction

- ✓ Shuffle memory contents to place all free memory together in one large block
- ✓ Compaction is possible *only* if relocation is dynamic, and is done at execution time
- ✓ I/O problem
  - Latch job in memory while it is involved in I/O
  - Do I/O only into OS buffers



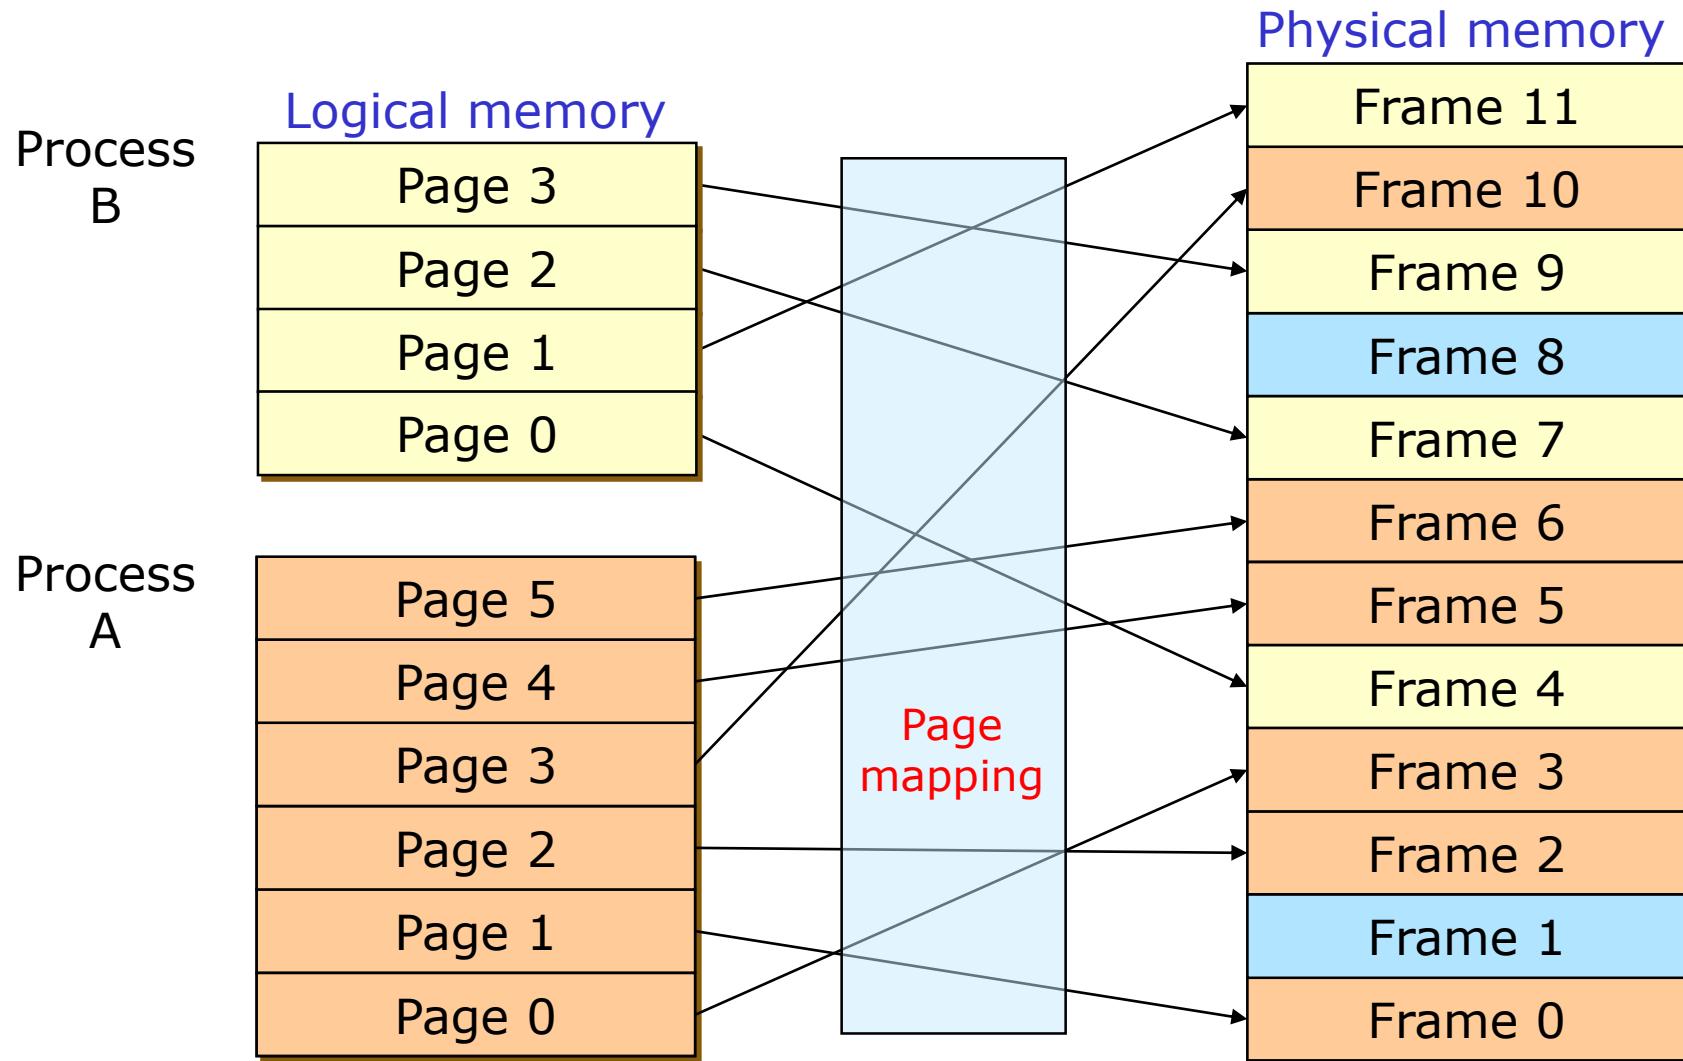
# Paging

---

- Physical address space of a process can be noncontiguous
  - ✓ Process is allocated physical memory whenever the latter is available
- Divide physical memory into fixed-sized blocks called **frames**
- Divide logical memory into blocks of same size called **pages**
  - ✓ Frame (or Page) size is power of 2 (typically, 512 bytes ~ 8192 bytes)
- Keep track of all free frames
- To run a program of size  $n$  pages, need to find  $n$  free frames and load program
- Set up a page table to translate logical to physical addresses
- Internal fragmentation



# Paging (Cont'd)



## ■ User's perspective

- ✓ Users (and processes) view memory as one contiguous address space from 0 through N
  - Logical (or virtual) address space
- ✓ In reality, pages are scattered throughout the physical memory
  - Logical-to-physical mapping
  - This mapping is invisible to the program
- ✓ Protection is provided because a program cannot reference memory outside of its logical address space
  - The logical address 0xdeadcafe maps to different physical addresses for different processes.



# **Address Translation Scheme**

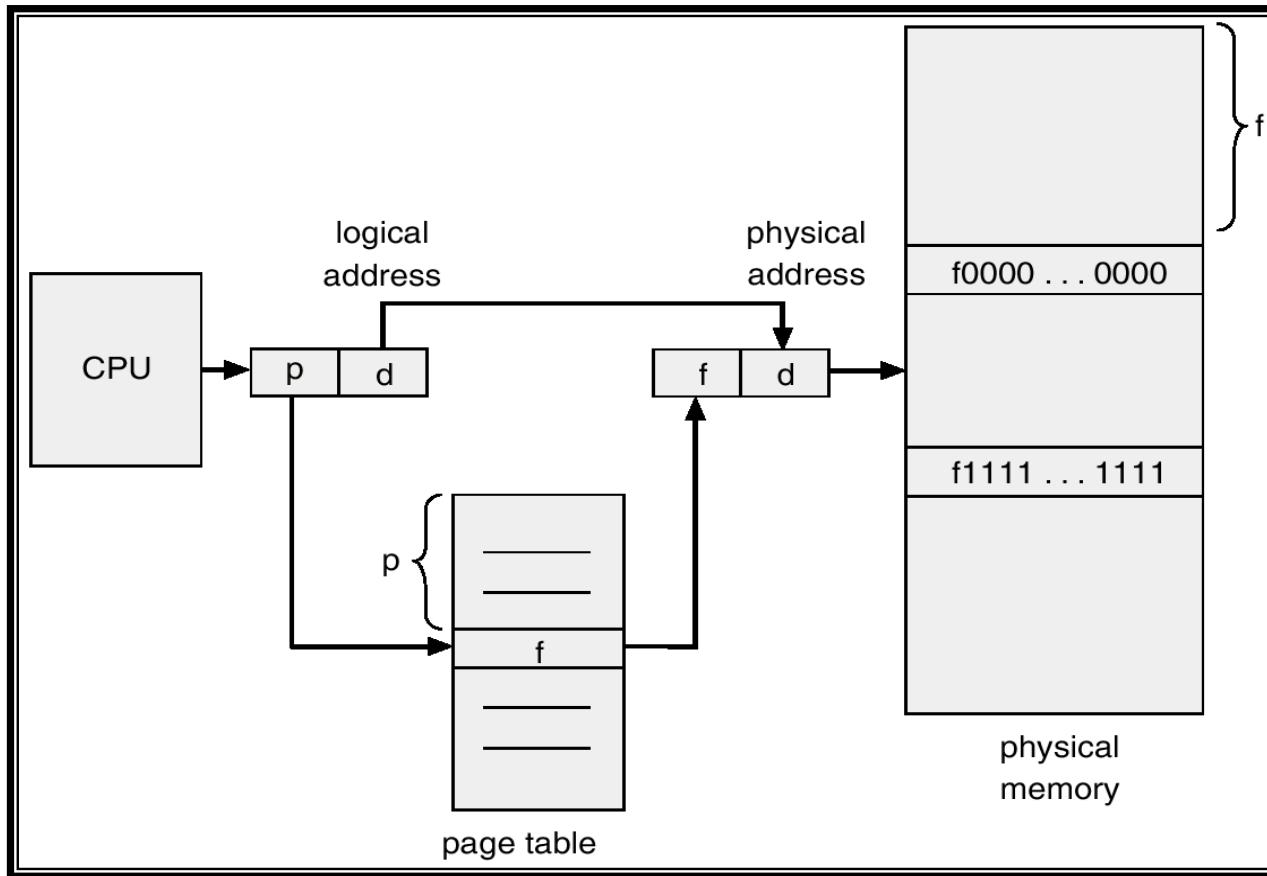
---

- Address generated by CPU is divided into:
  - ✓ *Page number (p)* – used as an index into a *page table* which contains base address of each page in physical memory
  - ✓ *Page offset (d)* – combined with base address to define the physical memory address that is sent to the memory unit



# *Address Translation Architecture*

---



## ■ Translating addresses

- ✓ A logical address has two parts:  
 $\langle \text{page number} :: \text{offset} \rangle$
- ✓ Page number is an index into a page table
- ✓ Page table determines frame number
- ✓ Physical address is  $\langle \text{frame number} :: \text{offset} \rangle$

## ■ Page tables

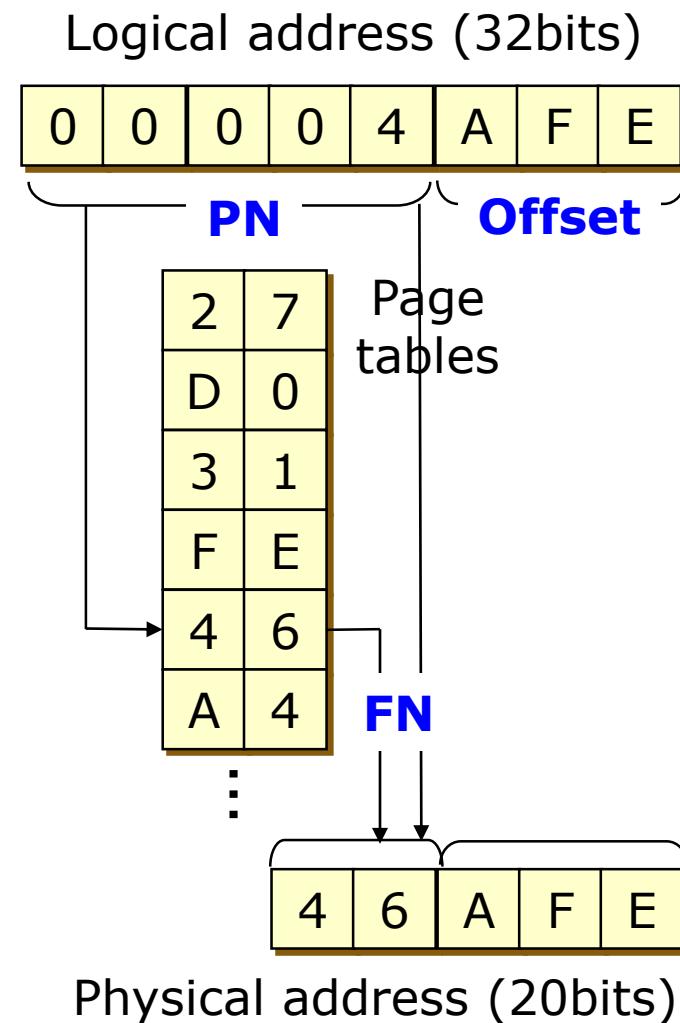
- ✓ Managed by OS
- ✓ Map page number to frame number
  - Page number is the index into the page table that determines frame number
- ✓ One page table entry per page in virtual address space

## ■ Cf) Frame table



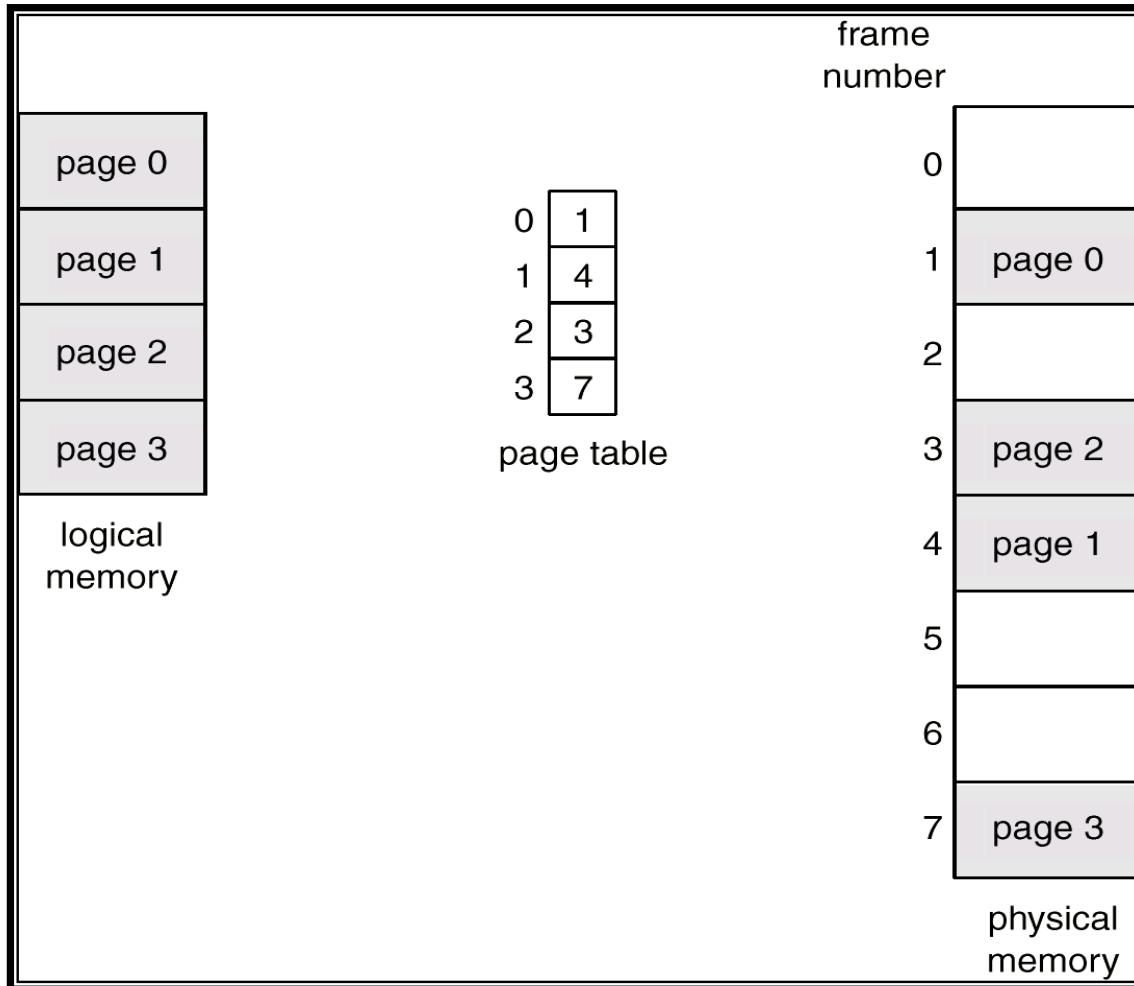
# Paging Example

- Logical address: 32 bits
- Physical address: 20 bits
- Page size: 4KB
- Offset: 12 bits
- Page Number: 20 bits
- Page table entries:  $2^{20}$

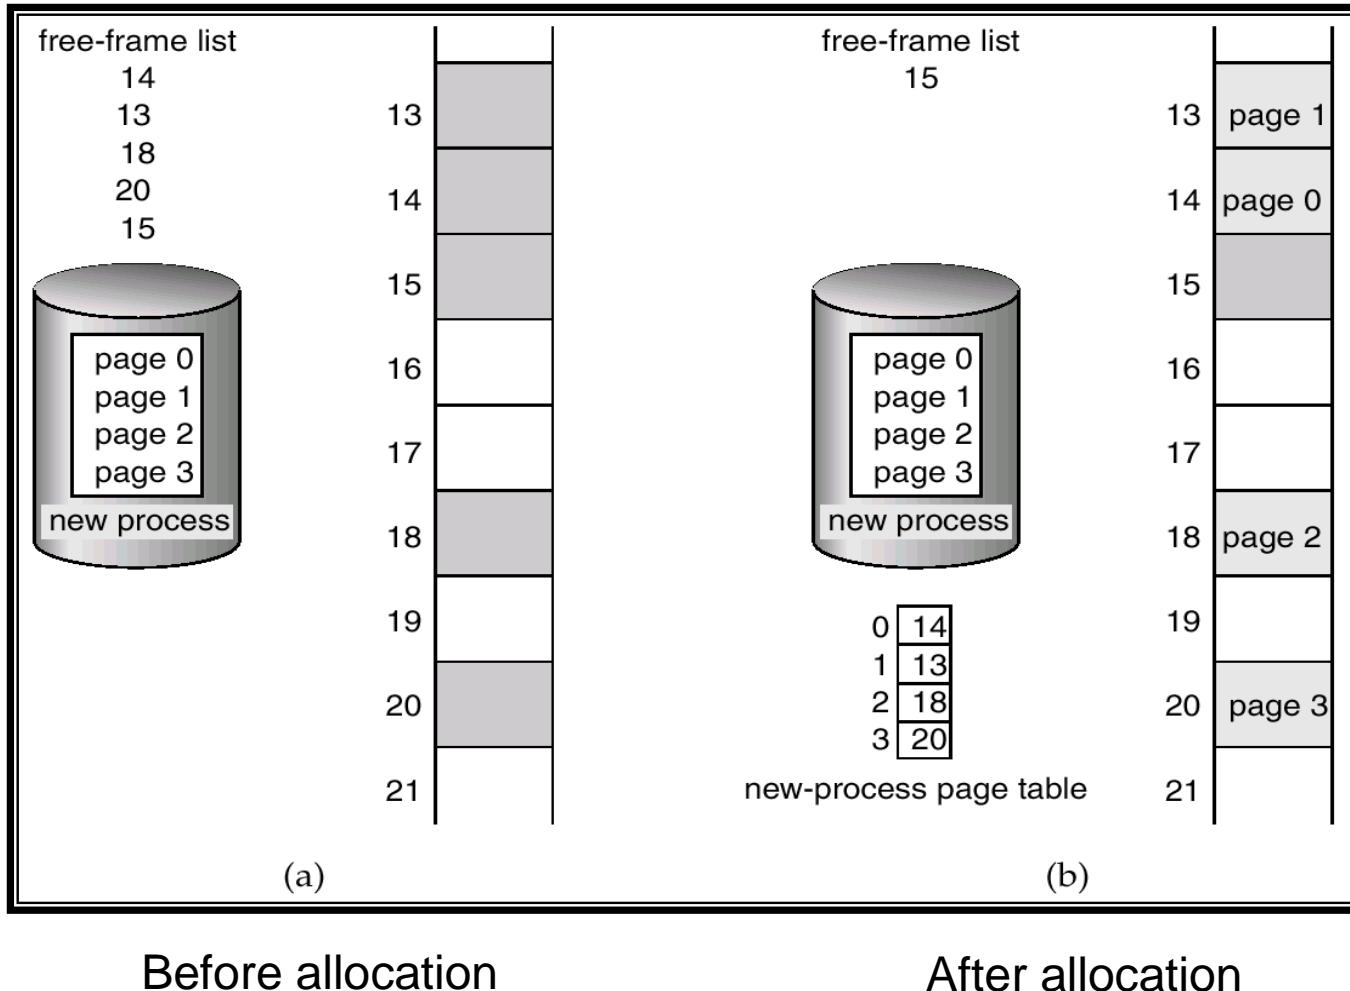


# Paging Example

---



# Free Frames



Before allocation

After allocation

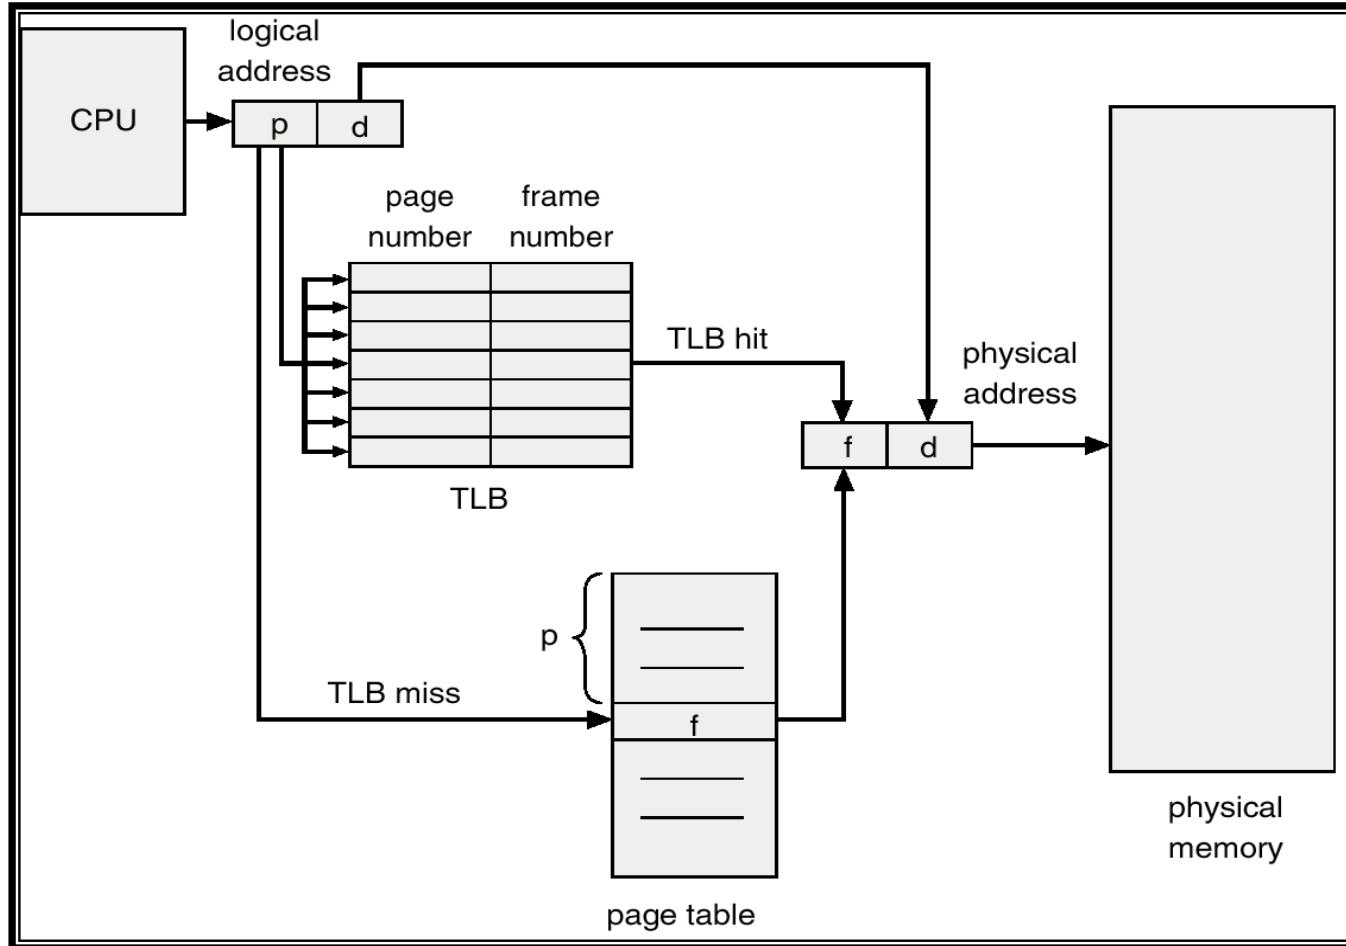
# **Implementation of Page Table**

---

- Page table is kept in main memory
- *Page-table base register (PTBR)* points to the page table
- *Page-table length register (PTLR)* indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
  - ✓ One for the page table and one for the data/instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called *associative memory* or *translation look-aside buffers (TLBs)*



# Paging Hardware With TLB



# Associative Memory

---

## ■ Associative memory – parallel search

| Page # | Frame # |
|--------|---------|
|        |         |
|        |         |
|        |         |
|        |         |

## ■ Address translation ( $A'$ , $A''$ )

- ✓ If  $A'$  is in associative register, get frame # out
- ✓ Otherwise get frame # from page table in memory

## ■ Cf) CAM (Content Addressable Memory)



## ■ Making address translation efficient

- ✓ Original page table scheme doubled the cost of memory lookups
  - One lookup into the page table, another to fetch the data
- ✓ Two-level page tables triple the cost!
  - Two lookups into the page tables, a third to fetch the data
  - And this assumes the page table is in memory
- ✓ How can we make this more efficient?
  - Goal: make fetching from a virtual address about as efficient as fetching from a physical address
  - Solutions:
    - Cache the virtual-to-physical translation in hardware
    - Translation Look-aside Buffer (TLB)
    - TLB managed by the Memory Management Unit (MMU)



## ■ Translation Look-aside Buffers

- ✓ Translate logical(virtual) page #s into page table entries (PTEs)  
(not physical address)
- ✓ Can be done in a single machine cycle

| <b>Valid</b> | <b>Virtual page</b> | <b>Modified</b> | <b>Protection</b> | <b>Page frame</b> |
|--------------|---------------------|-----------------|-------------------|-------------------|
| 1            | 140                 | 1               | RW                | 31                |
| 1            | 20                  | 0               | R X               | 38                |
| 1            | 130                 | 1               | RW                | 29                |
| 1            | 129                 | 1               | RW                | 62                |
| 1            | 19                  | 0               | R X               | 50                |
| 1            | 21                  | 0               | R X               | 45                |
| 1            | 860                 | 1               | RW                | 14                |
| 1            | 861                 | 1               | RW                | 75                |

## ■ TLB is implemented in hardware

- ✓ Fully associative cache (all entries looked up in parallel)
- ✓ Cache tags are logical(virtual) page numbers
- ✓ Cache values are PTEs (entries from page tables)
- ✓ With PTE+offset, MMU can directly calculate the physical address

## ■ TLBs exploit locality

- ✓ Processes only use a handful of pages at a time
  - 16~48 entries in TLB is typical (64~192KB)
  - Can hold the “hot set” or “working set” of process
- ✓ Hit rates are therefore really important

## ■ Locality

- ✓ Temporal locality vs. Spatial locality



## ■ Handling TLB misses

- ✓ Address translations are mostly handled by the TLB
  - > 99% of translations, but there are TLB misses occasionally
  - In case of a miss, who places translations into the TLB?
- ✓ Hardware (MMU): Intel x86
  - Knows where page tables are in memory
  - OS maintains tables, HW access them directly
  - Page tables have to be in hardware-defined format
- ✓ Software loaded TLB (OS)
  - TLB miss faults to OS, OS finds right PTE and loads TLB
  - Must be fast (but, 20-200 cycles typically)
  - CPU ISA has instructions for TLB manipulation
  - Page tables can be in any format convenient for OS (flexible)



## ■ Managing TLBs

- ✓ OS ensures that TLB and page tables are consistent
  - When OS changes the protection bits of a PTE, it needs to invalidate the PTE if it is in the TLB
- ✓ Reload TLB on a process context switch
  - Remember, each process typically has its own page tables
  - Need to invalidate all the entries in TLB (flush TLB)
  - In IA32, TLB is flushed automatically when the contents of CR3 (page directory base register) is changed
  - Cf) Alternatively, we can store the PID as part of the TLB entry, but this is expensive
- ✓ When the TLB misses, and a new PTE is loaded, a cached PTE must be evicted
  - Choosing a victim PTE called the “TLB replacement policy”
  - Implemented in hardware, usually simple (e.g., LRU)



# *Effective Access Time*

---

- Associative Lookup =  $\varepsilon$  time unit
- Assume memory cycle time is 1 microsecond
- Hit ratio
  - ✓ Percentage of times that a page number is found in the associative registers
  - ✓ Ratio related to number of associative registers
- Hit ratio =  $\alpha$
- Effective Access Time (EAT)
$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$



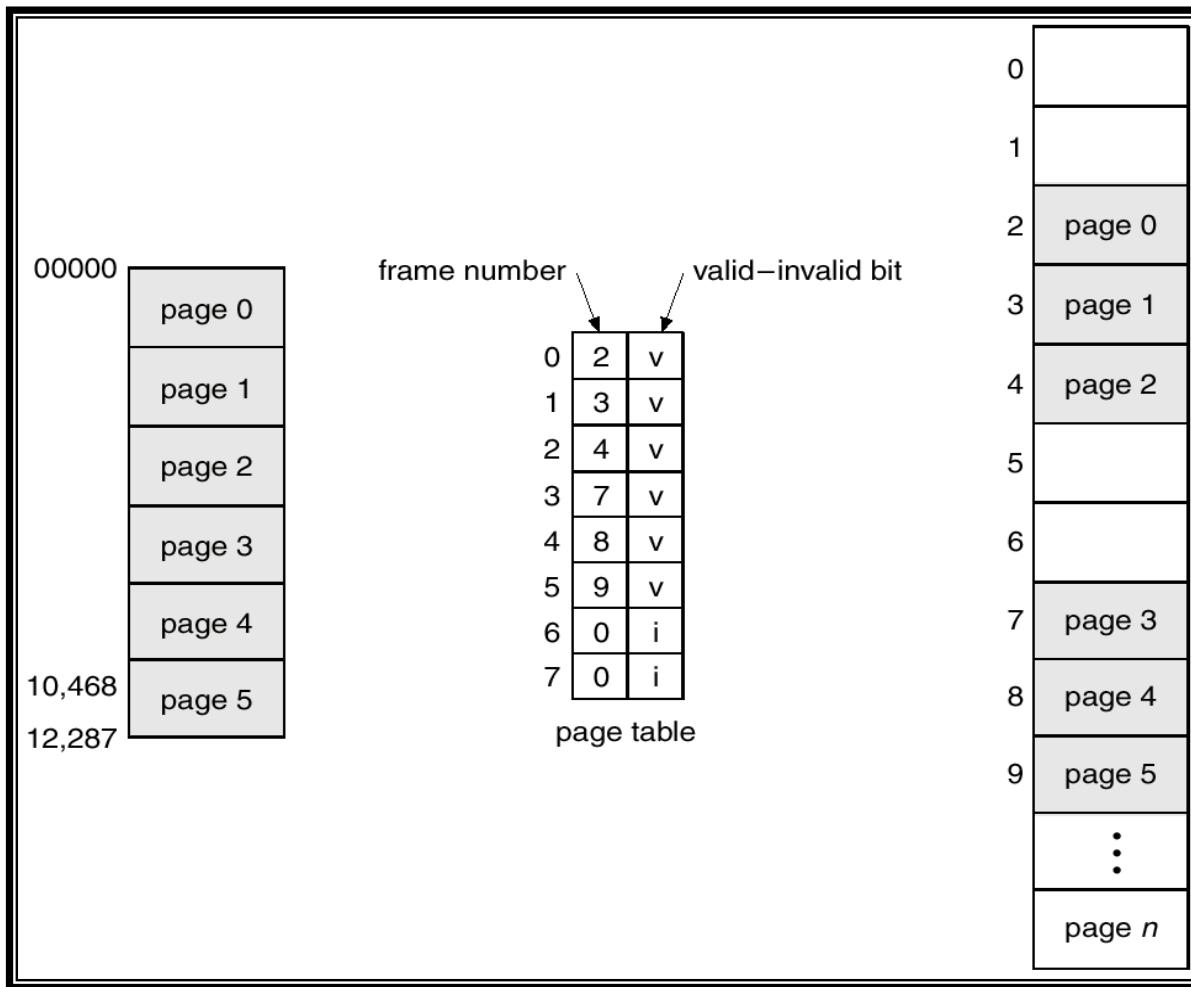
# Memory Protection

---

- Memory protection implemented by associating protection bit with each frame
- *Valid-invalid* bit attached to each entry in the page table:
  - ✓ “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - ✓ “invalid” indicates that the page is not in the process’ logical address space
- Finer level of protection is possible for valid pages
  - ✓ Provide read-only, read-write, or execute-only protection



# *Valid (v) or Invalid (i) Bit In A Page Table*



# *Page Table Entries (PTEs)*

---



- Valid bit (V) says whether or not the PTE can be used
  - ✓ It is checked each time a virtual address is used
- Reference bit (R) says whether the page has been accessed
  - ✓ It is set when a read or write to the page occurs
- Modify bit (M) says whether or not the page is dirty
  - ✓ It is set when a write to the page occurs
- Protection bits (Prot) control which operations are allowed on the page
  - ✓ Read, Write, Execute, etc.
- Frame number (FN) determines physical page

# *Page Table Structure*

---

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables



## ■ Managing page tables

- ✓ Space overhead of page tables
  - The size of the page table for a 32-bit address space with 4KB pages = 4MB (per process)
- ✓ How can we reduce this overhead?
  - Observation: Only need to map the portion of the address space actually being used (tiny fraction of entire address space)
- ✓ How do we only map what is being used?
  - Make the page table structure dynamically extensible
  - Use another level of indirection:
    - Two-level, hierarchical, hashed, etc.



# *Hierarchical Page Tables*

---

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table



# Two-Level Paging Example

---

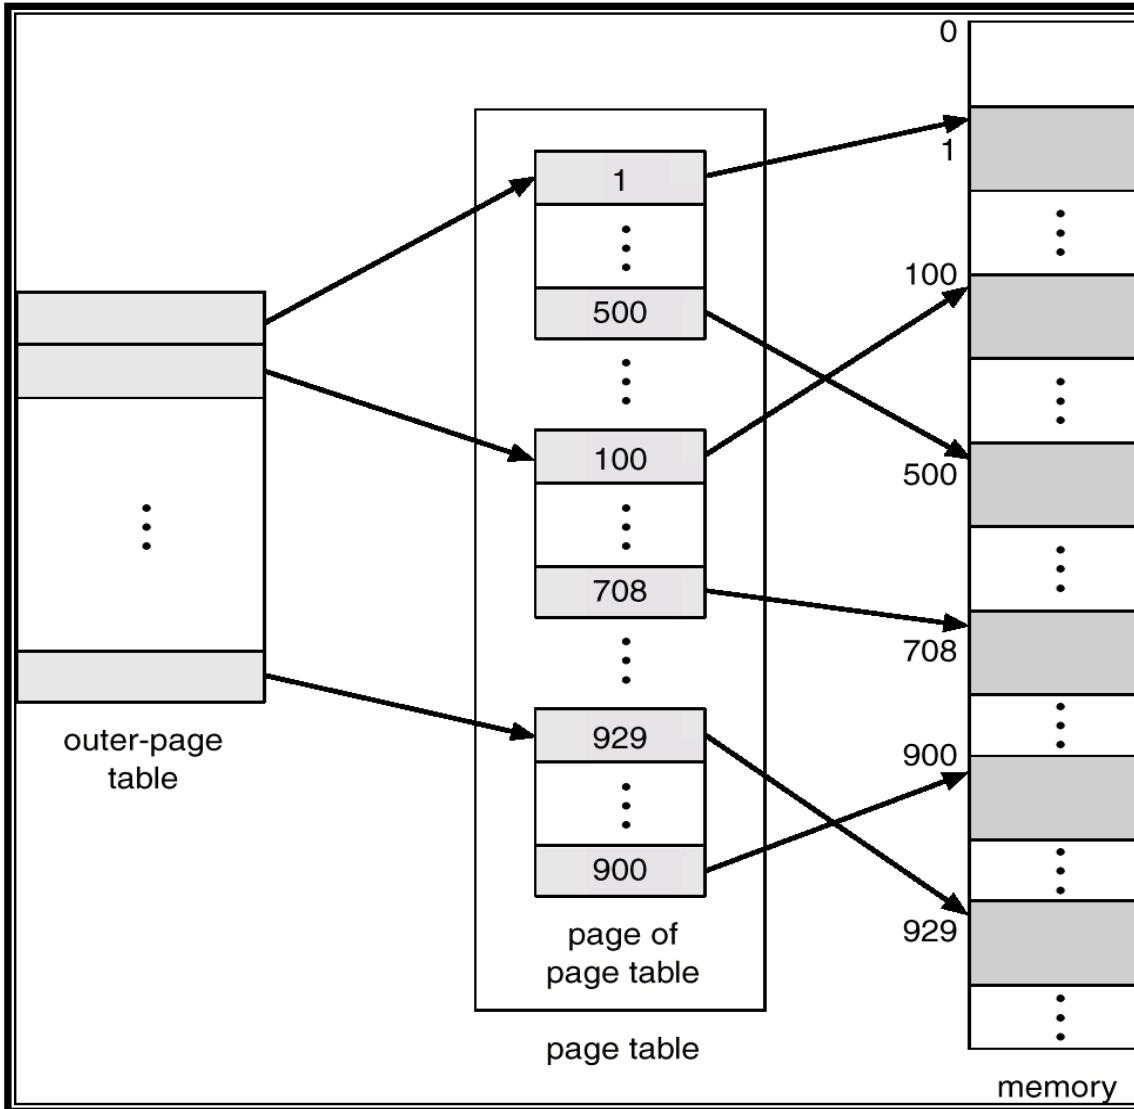
- A logical address (on 32-bit machine with 4K page size) is divided into:
  - ✓ a page number consisting of 20 bits
  - ✓ a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
  - ✓ a 10-bit page number
  - ✓ a 10-bit page offset
- Thus, a logical address is as follows:

| page number |       | page offset |
|-------------|-------|-------------|
| $p_1$       | $p_2$ | $d$         |
| 10          | 10    | 12          |

- ✓ where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the outer page table

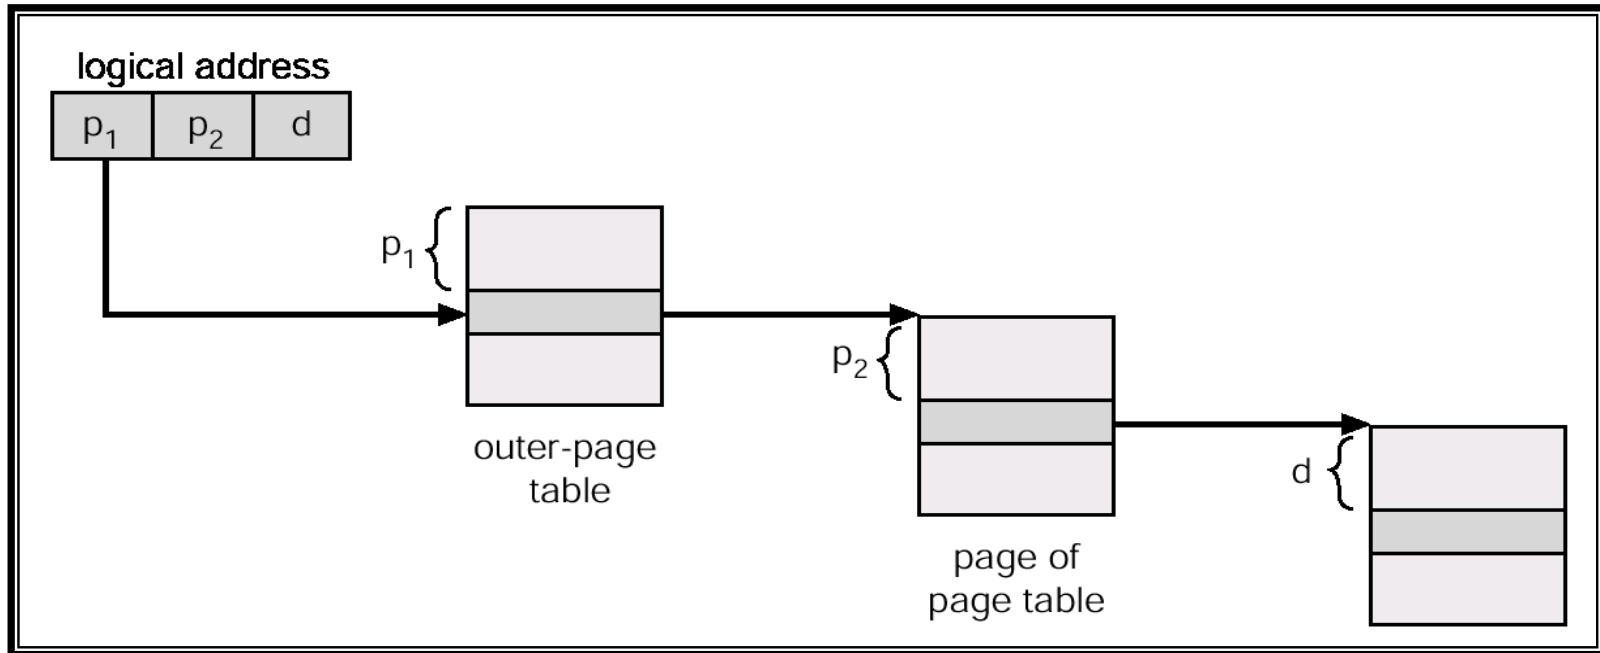


# Two-Level Page-Table Scheme



# **Address-Translation Scheme**

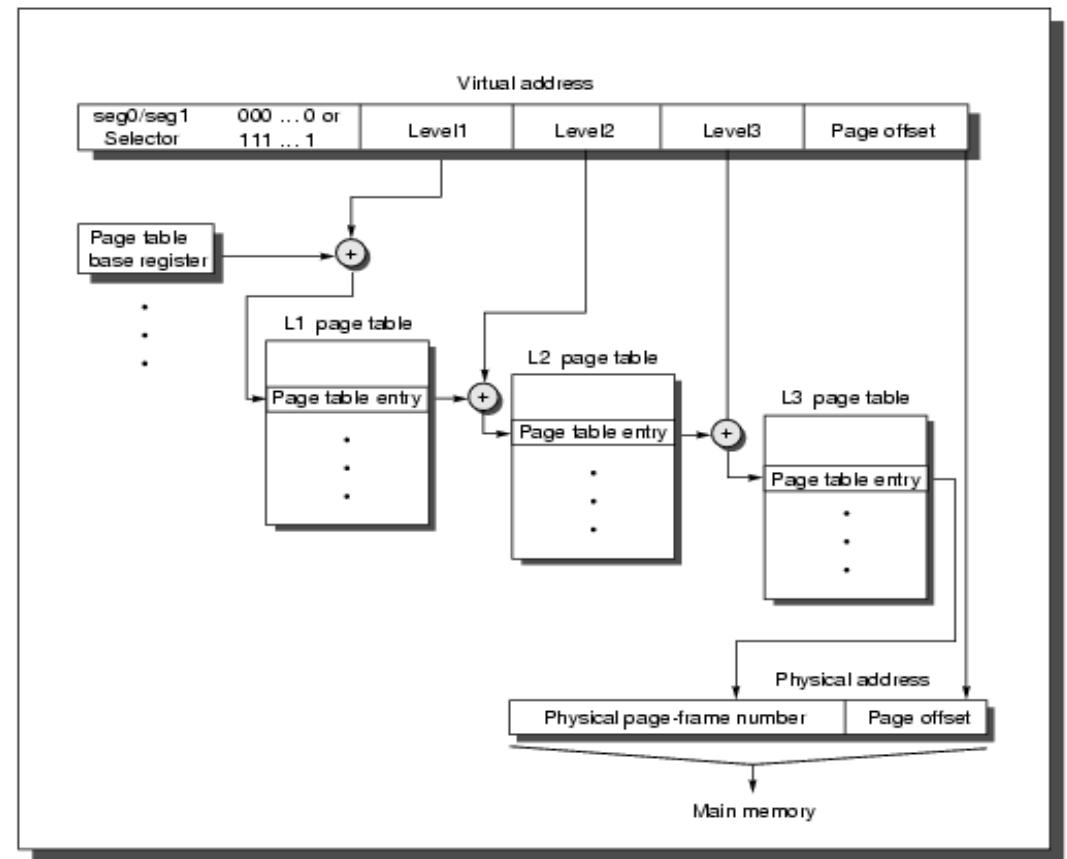
- Address-translation scheme for a two-level 32-bit paging architecture



# Multi-level Page Tables

## Address translation in Alpha AXP Architecture

- ✓ Three-level page tables
- ✓ 64-bit address divided into 3 segments (coded in bits63/62)
  - seg0 (0x): user code
  - seg1 (11): user stack
  - kseg (10): kernel
- ✓ Alpha 21064
  - Page size: 8KB
  - Virtual address: 43bits
  - Each page table is one page long



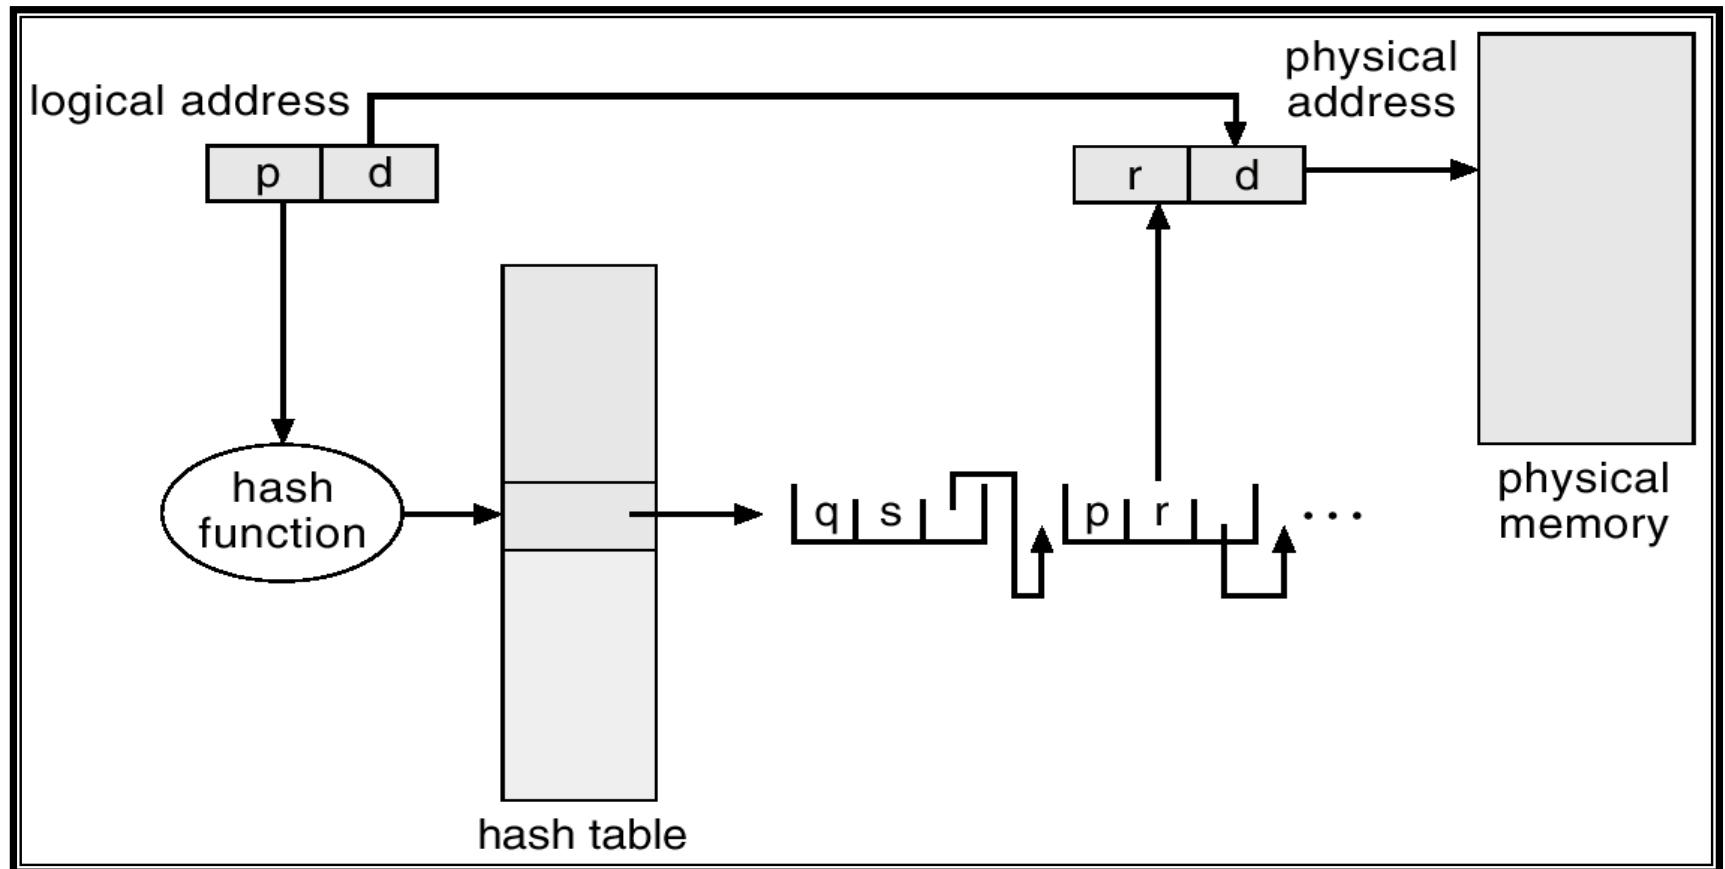
# *Hashed Page Tables*

---

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
  - ✓ This page table contains a chain of elements hashing to the same location
- Virtual page numbers are compared in this chain searching for a match
  - ✓ If a match is found, the corresponding physical frame is extracted



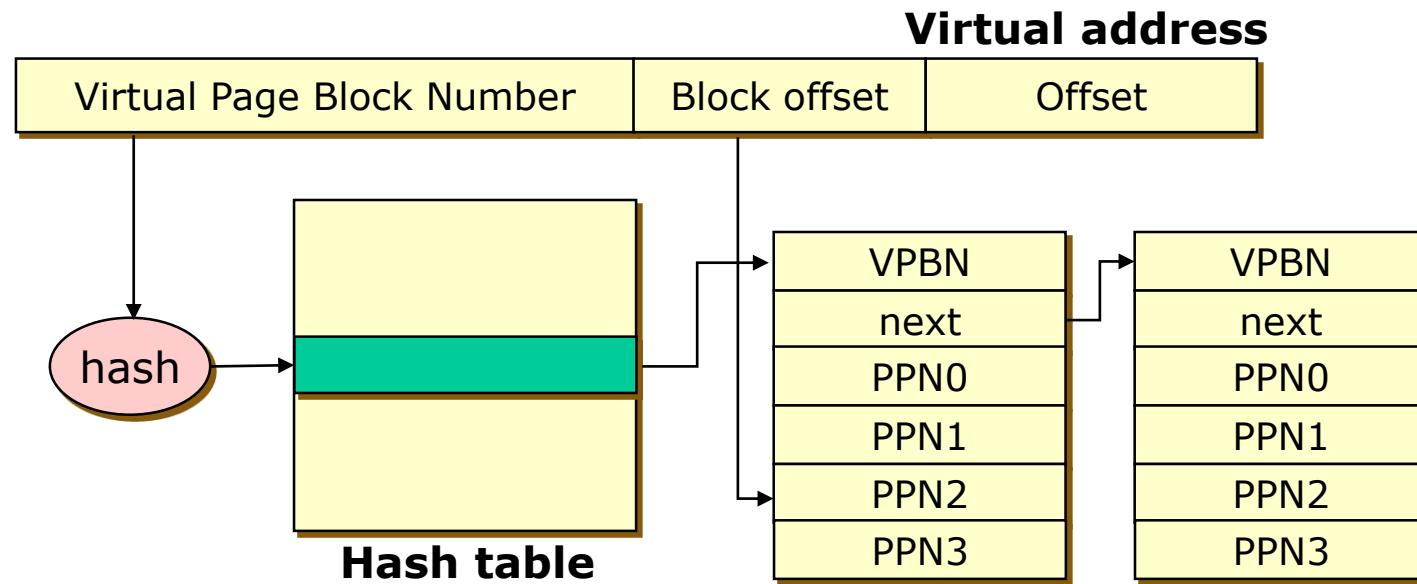
# Hashed Page Table



# Hashed Page Tables

## Clustered page tables

- ✓ A variant of hash page tables with the difference that each entry stores mapping information for a block of consecutive page tables



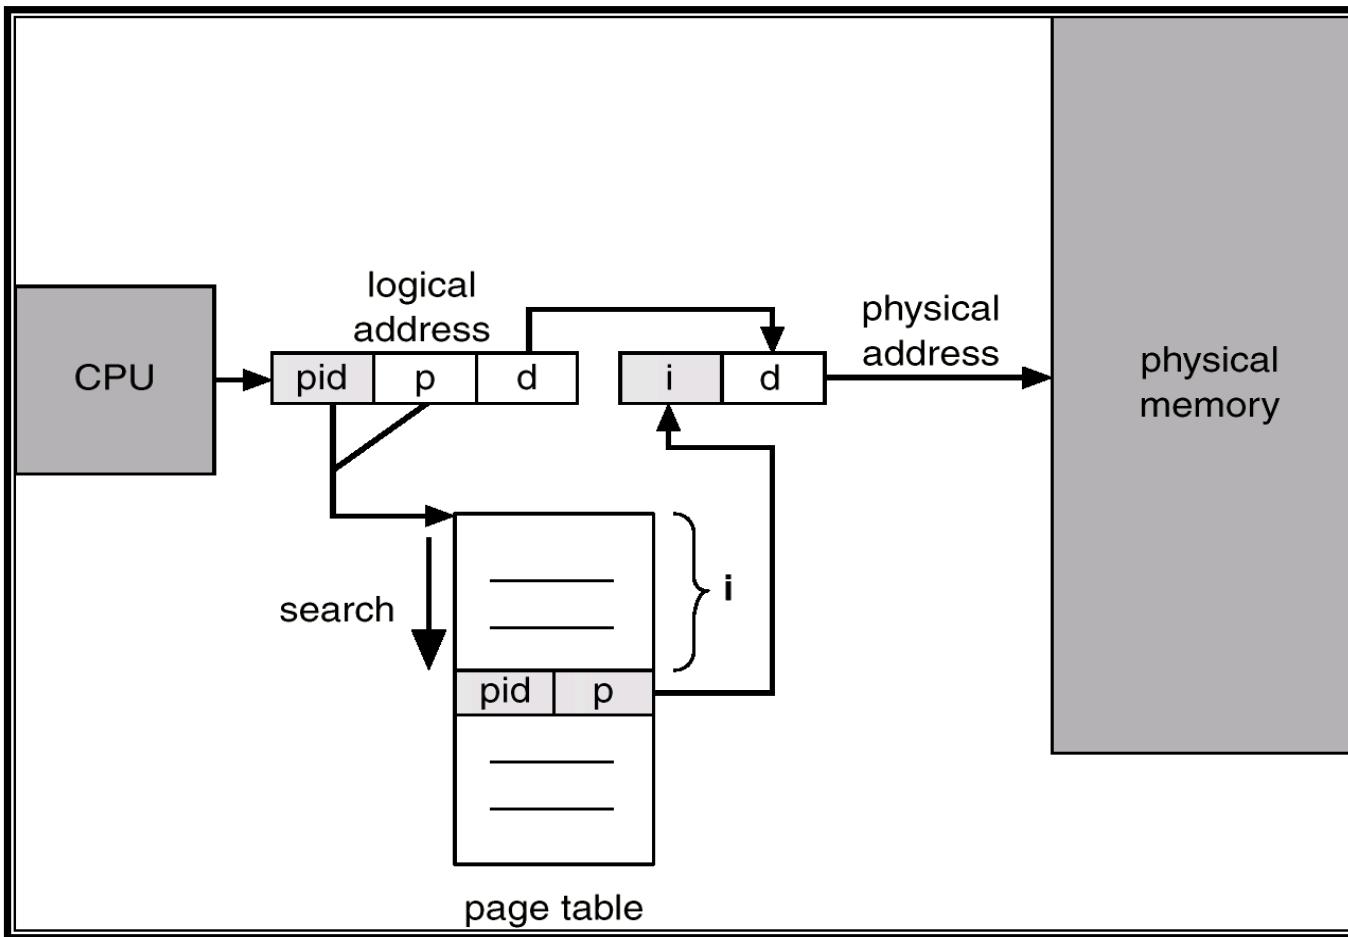
# *Inverted Page Table*

---

- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one, or at most a few, page-table entries



# Inverted Page Table Architecture



# **Shared Pages**

---

## ■ Shared code

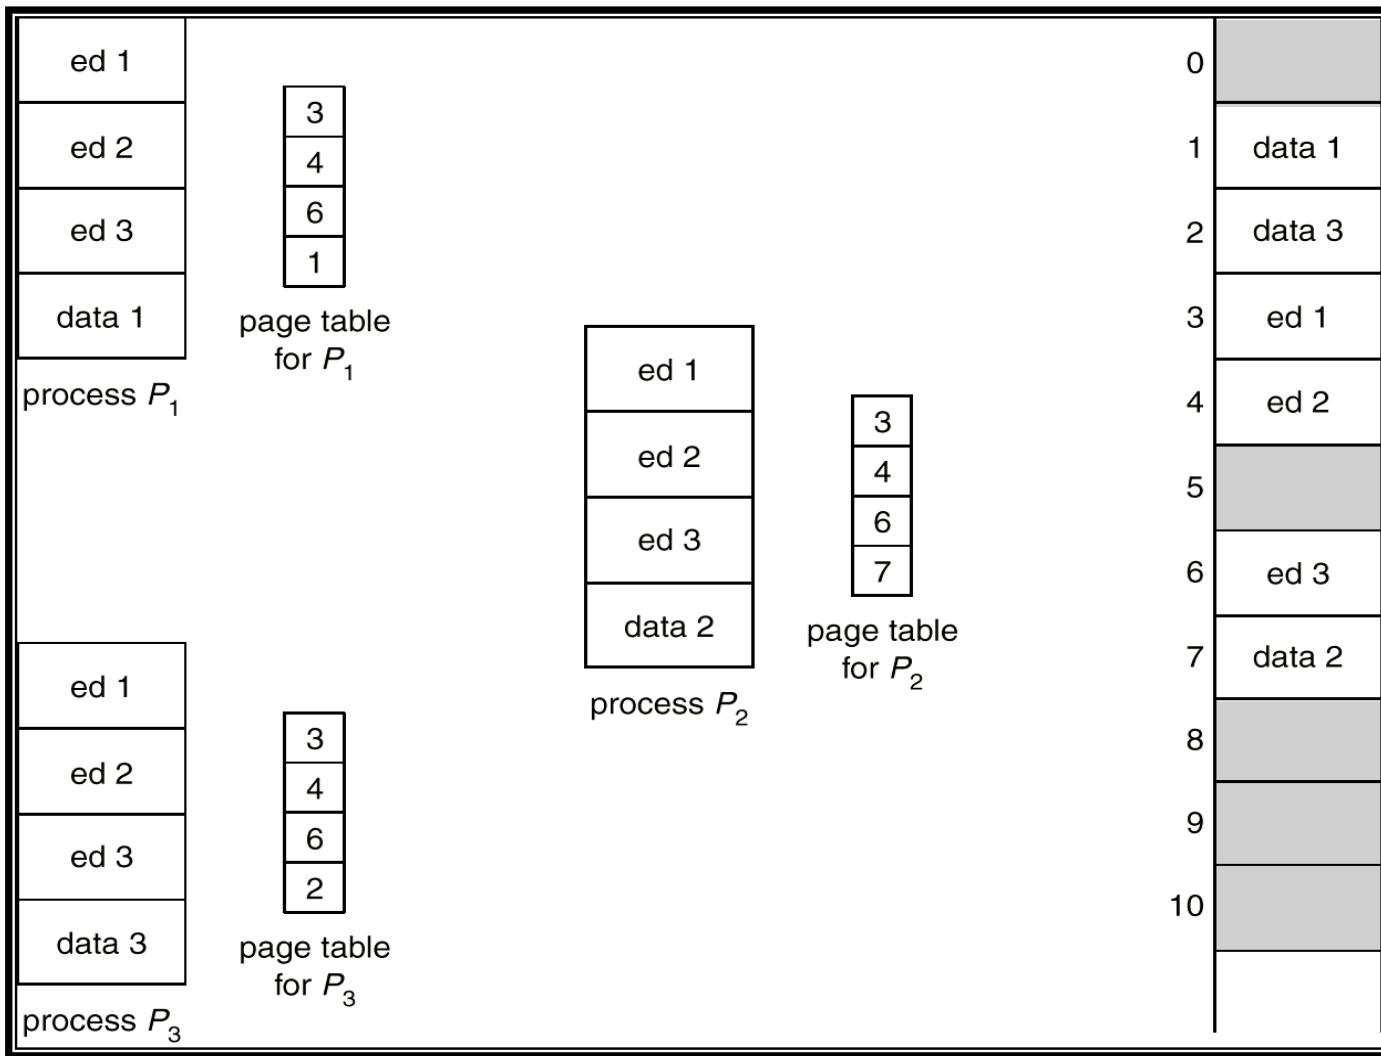
- ✓ One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems)
- ✓ Shared code must appear in same location in the logical address space of all processes

## ■ Private code and data

- ✓ Each process keeps a separate copy of the code and data
- ✓ The pages for the private code and data can appear anywhere in the logical address space



# Shared Pages Example



# *Advantages of Paging*

---

- Easy to allocate physical memory
  - ✓ Physical memory is allocated from free list of frames
  - ✓ To allocate a frame, just remove it from its free list
- No external fragmentation
- Easy to “page out” chunks of a program
  - ✓ All chunks are the same size (page size)
  - ✓ Use valid bit to detect reference to “paged-out” pages
  - ✓ Pages sizes are usually chosen to be convenient multiple of disk block sizes
- Easy to protect pages from illegal accesses
- Easy to share pages



# *Disadvantages of Paging*

---

- Can still have internal fragmentation

- ✓ Process may not use memory in exact multiple of pages

- Memory reference overhead

- ✓ 2 references per address lookup (page table, then memory)
  - ✓ Solution
    - get a hardware support (TLB)

- Memory required to hold page tables can be large

- ✓ Need one page table entry(PTE) per page in virtual address space
  - ✓ 32-bit address space with 4KB pages =  $2^{20}$  PTEs
  - ✓ 4 bytes/PTE = 4MB per page table
  - ✓ OS's typically have separate page tables per process  
(25 processes = 100MB of page tables)
  - ✓ Solution
    - page the page tables, multi-level page tables, hashed page tables, inverted page tables, etc.



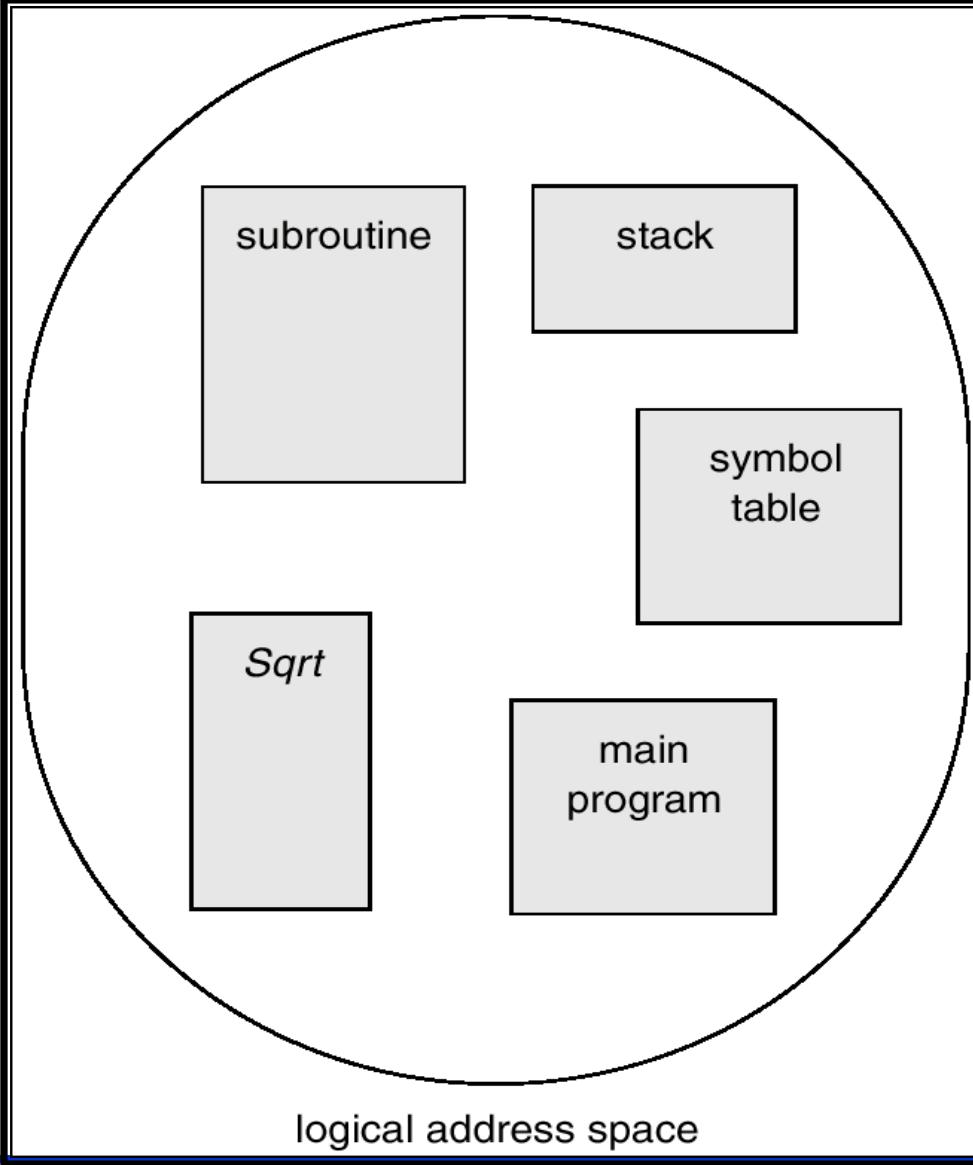
# **Segmentation**

---

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
- A segment is a logical unit such as:
  - ✓ Main program
  - ✓ Procedure
  - ✓ Function
  - ✓ Method
  - ✓ Object
  - ✓ Local variables, Global variables
  - ✓ Common block
  - ✓ Stack
  - ✓ Symbol table
  - ✓ Arrays

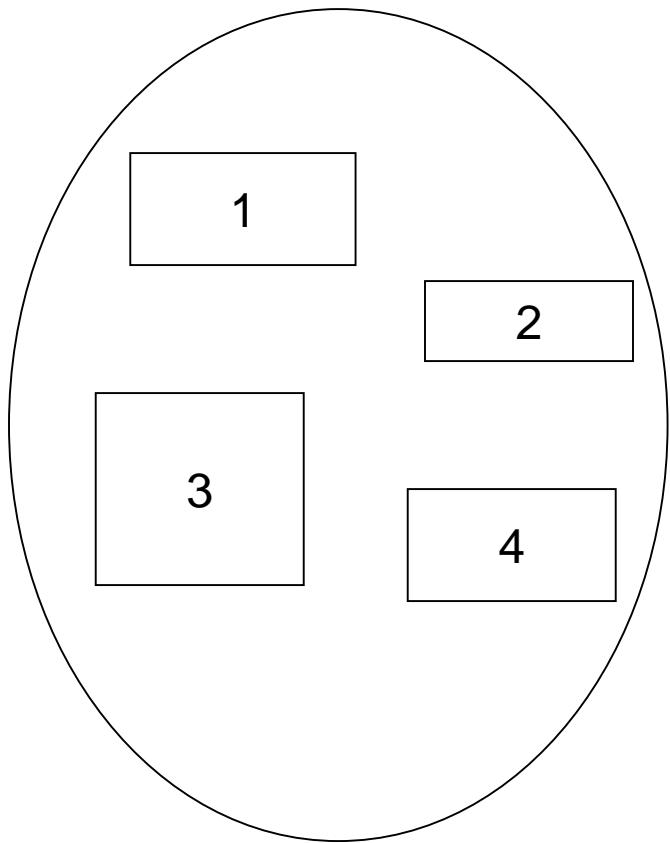


# User's View of a Program

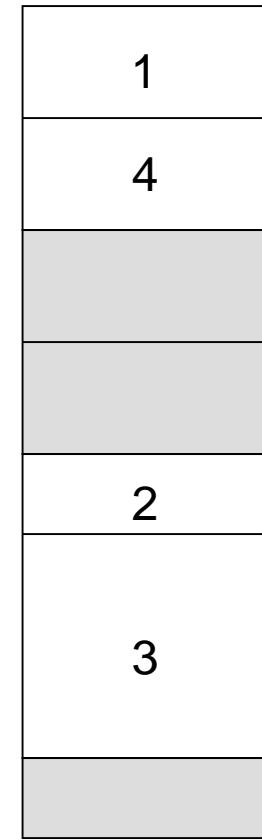


# *Logical View of Segmentation*

---



user space



physical memory space

# Segmentation Architecture

---

- Logical address consists of a two tuple:  
 $\langle \text{segment-number}, \text{offset} \rangle$
- *Segment table* – maps two-dimensional physical addresses
- Each table entry has:
  - ✓ base – contains the starting physical address where the segments reside in memory
  - ✓ *limit* – specifies the length of the segment
- *Segment-table base register (STBR)* points to the segment table's location in memory
- *Segment-table length register (STLR)* indicates number of segments used by a program
  - ✓ Segment number  $s$  is legal if  $s < \text{STLR}$



# ***Segmentation Architecture (Cont'd)***

---

## ■ Relocation

- ✓ dynamic
- ✓ by segment table

## ■ Sharing

- ✓ shared segments
- ✓ same segment number

## ■ Allocation

- ✓ first fit/best fit
- ✓ external fragmentation



# **Segmentation Architecture (Cont'd)**

---

## ■ Protection

- ✓ With each entry in segment table associate:
  - validation bit = 0  $\Rightarrow$  illegal segment
  - read/write/execute privileges

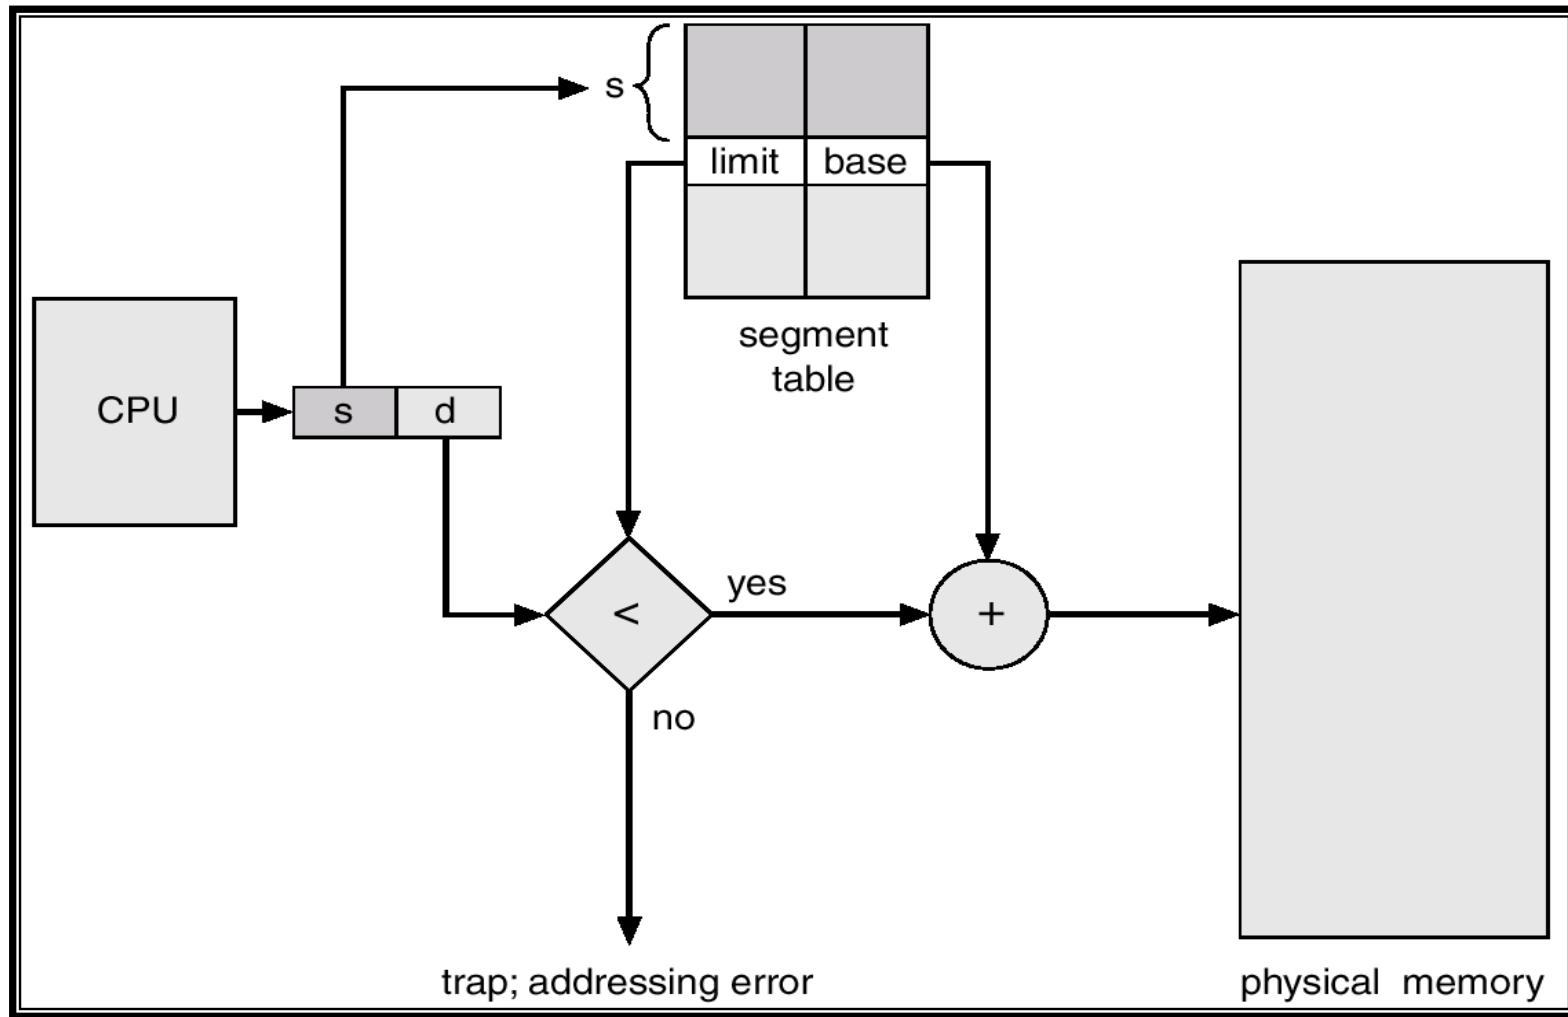
## ■ Protection bits associated with segments

- ✓ Code sharing occurs at segment level

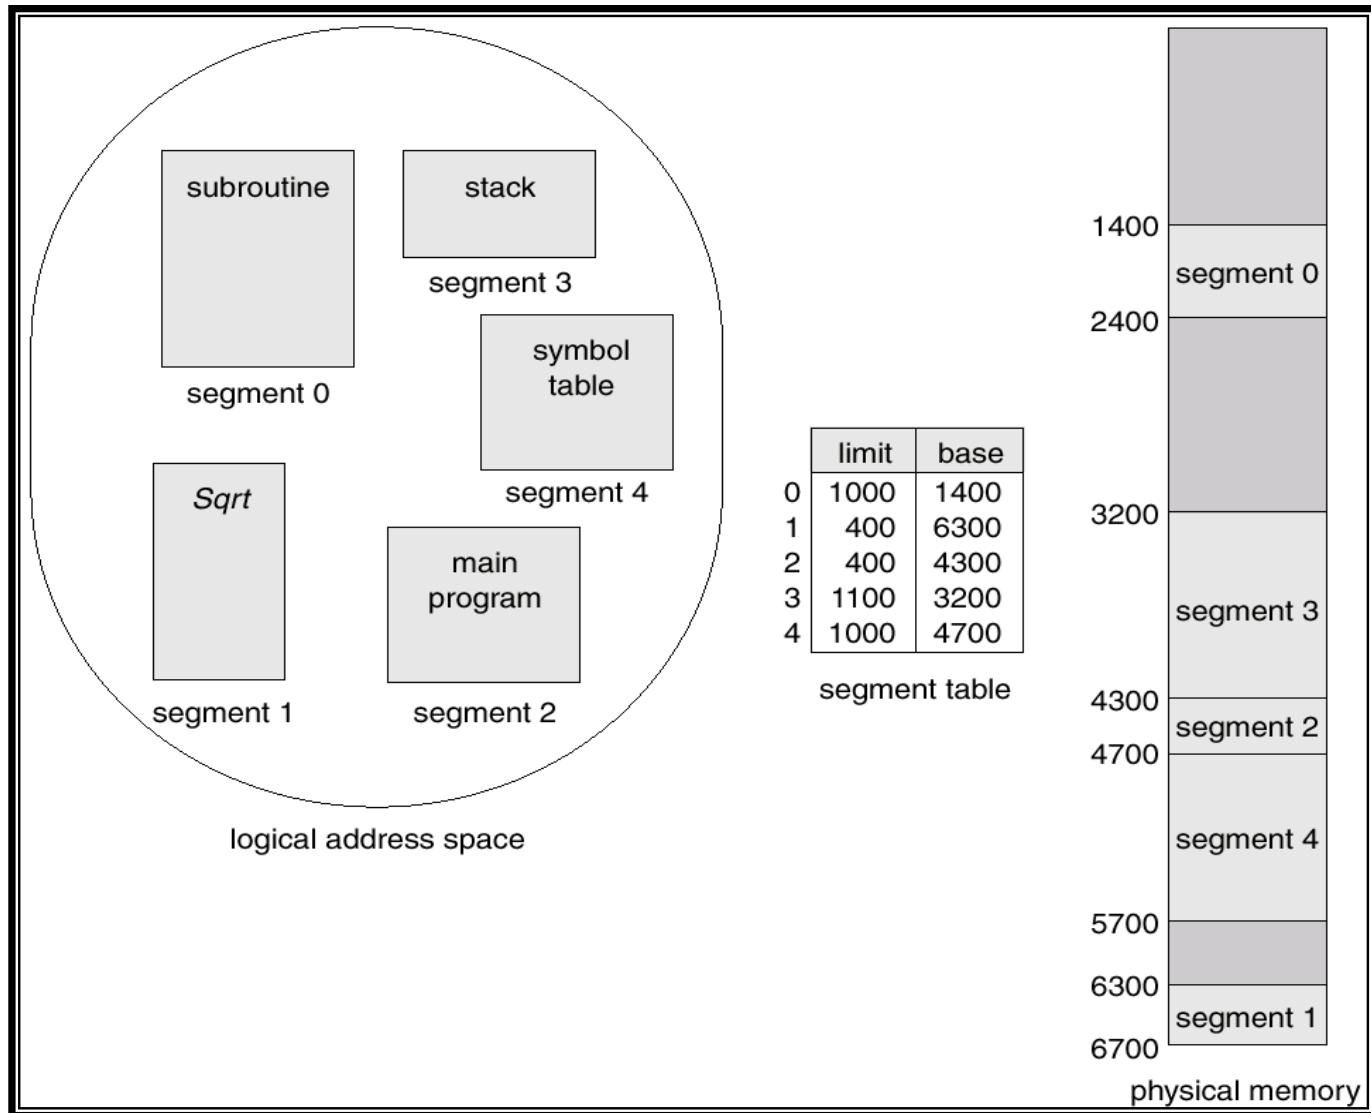
## ■ Since segments vary in length, memory allocation is a dynamic storage-allocation problem



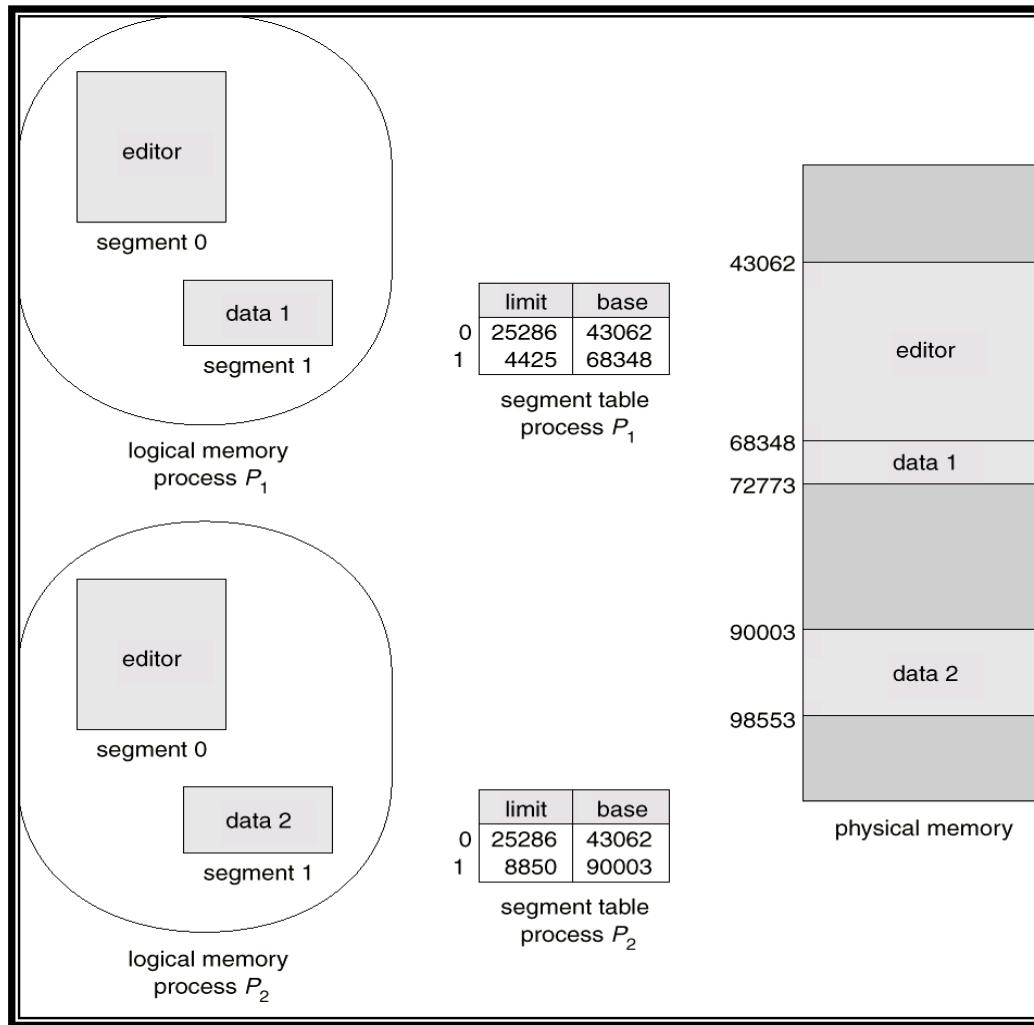
# *Segmentation Hardware*



# Example of Segmentation



# Sharing of Segments



# *Advantages of Segmentation*

---

- Simplifies the handling of data structures that are growing or shrinking
- Easy to protect segments
  - ✓ With each entry in segment table, associate a valid bit
  - ✓ Protection bits (read/write/execute) are also associated with each segment table entry
- Easy to share segments
  - ✓ Put same translation into base/limit pair
  - ✓ Code/data sharing occurs at segment level
  - ✓ e.g. shared libraries
- No internal fragmentation



# *Disadvantages of Segmentation*

---

## ■ Cross-segment addresses

- ✓ Segments need to have same segment # for pointers to them to be shared among processes
- ✓ Otherwise, use indirect addressing only

## ■ Large segment tables

- ✓ Keep in main memory, use hardware cache for speed

## ■ External fragmentation

- ✓ Since segments vary in length, memory allocation is a dynamic storage-allocation problem



# *Paging vs. Segmentation*

|                                 | <b>Paging</b>                          | <b>Segmentation</b>                                |
|---------------------------------|--|--|
| Block size                      | Fixed (4KB to 64KB)                    | Variable   |
| Linear address space            | 1                                      | Many   |
| Memory addressing               | One word<br>(page number + offset)     | Two words<br>(segment & offset)                    |
| Replacement                     | Easy<br>(all same size)                | Difficult<br>(find where segment fits)             |
| Fragmentation                   | Internal                               | External   |
| Disk traffic                    | Efficient<br>(optimized for page size) | Inefficient<br>(may have small or large transfers) |
| Transparent to the programmers? | Yes                                    | No   |



# *Paging vs. Segmentation (Cont'd)*

|   | <b>Paging</b>  | <b>Segmentation</b>  |
|---|--|--|
| Can the total address space exceed the size of physical memory? | Yes  | Yes  |
| Can codes and data be distinguished and separately protected?   | No   | Yes  |
| Can tables whose size fluctuates be accommodated easily?        | No   | Yes  |
| Is sharing of codes between users facilitated?                  | No   | Yes  |
| Why was this technique invented?                                | To get a large linear address space without having to buy more physical memory | To allow programs and data to be broken up into logically independent address spaces and to aid sharing and protection |



# *Paging vs. Segmentation (Cont'd)*

---

## ■ Hybrid approaches

- ✓ Paged segments
  - Segmentation with Paging
  - Segments are a multiple of a page size
- ✓ Multiple page sizes
  - 4KB, 2MB, and 4MB page sizes are supported in IA32
  - 8KB, 16KB, 32KB or 64KB in Alpha AXP Architecture  
(43, 47, 51, or 55 bits virtual address)



# ***Segmentation with Paging***

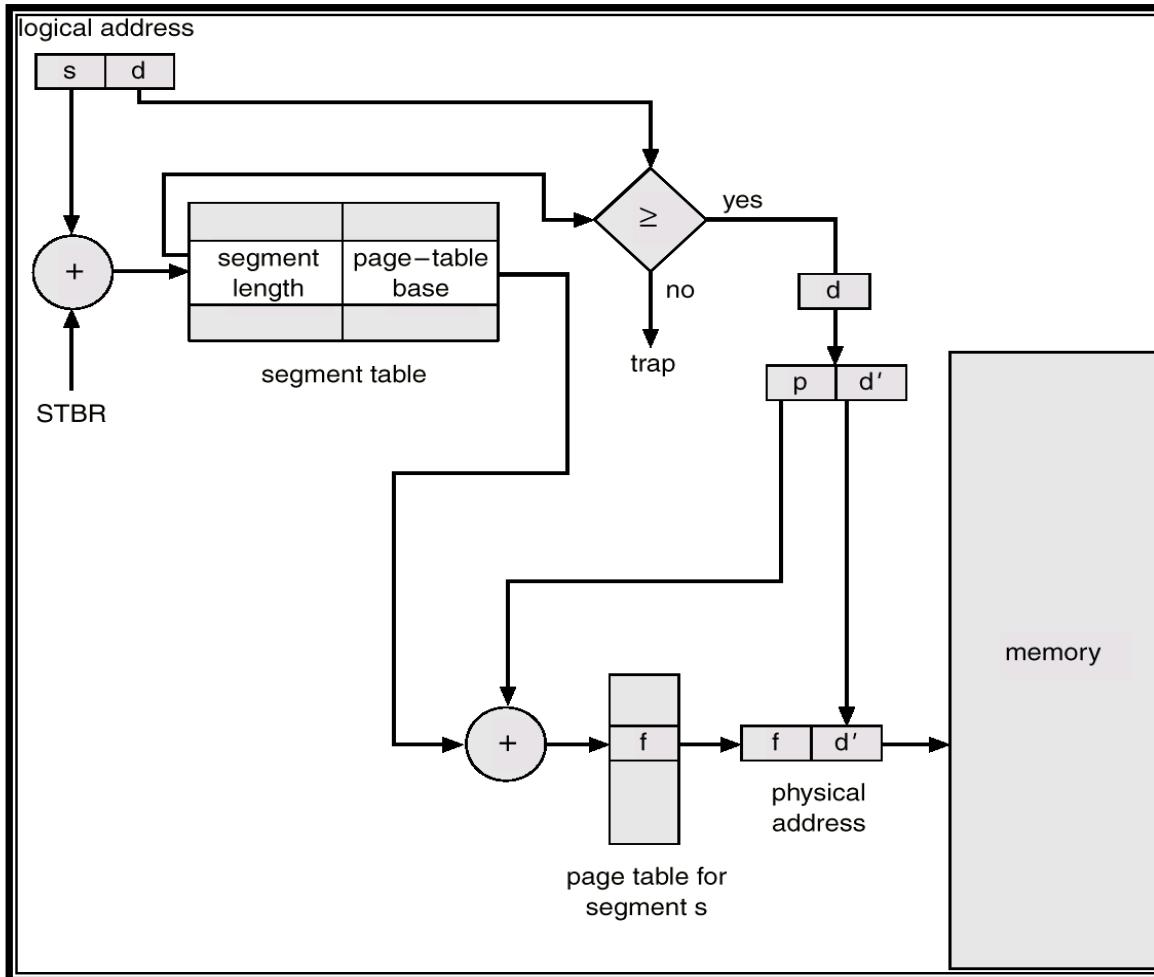
---

## ■ Combine segmentation and paging

- ✓ Use segments to manage logically related units
  - Code, data, heap, etc.
  - Segments vary in size, but usually large (multiple pages)
- ✓ Use pages to partition segments into fixed size chunks
  - Makes segments easier to manage with in physical memory
  - Segments become “pageable” – rather than moving segments into and out of memory, just move page portions of segments
  - Need to allocate page table entries only for those pieces of the segments that have themselves been allocated
  - No external fragmentation
- ✓ The IA32 supports segments and paging

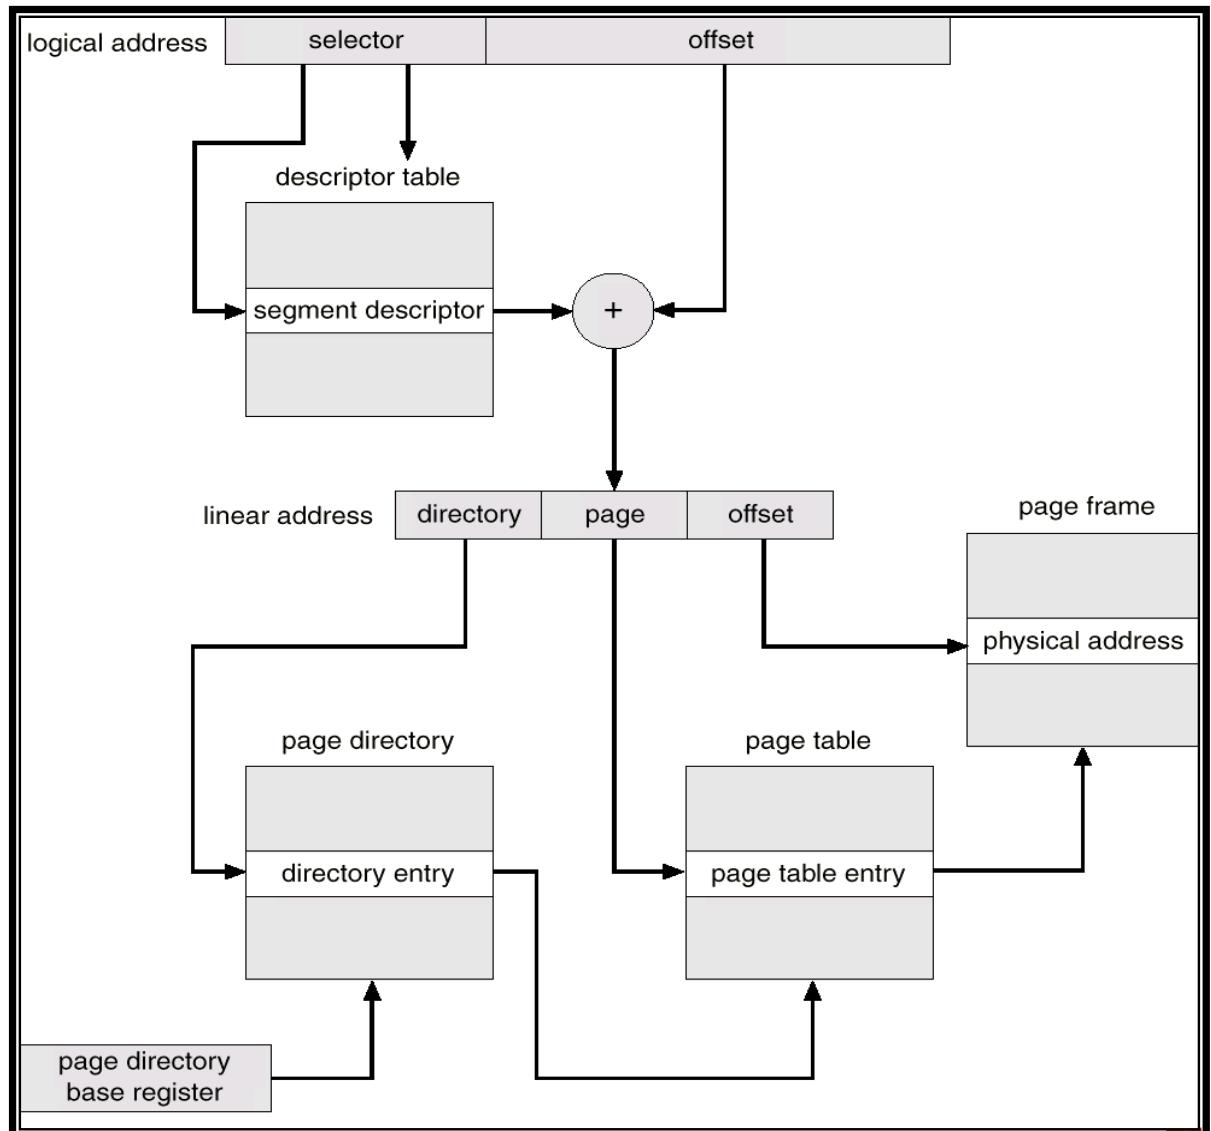


# MULTICS Address Translation Scheme



# *Segmentation with Paging – Intel Pentium*

- The Intel Pentium uses segmentation with paging for memory management with a two-level paging scheme



## ■ Goals of memory management

- ✓ To provide a convenient abstraction for programming
  - Logical address vs. physical address
- ✓ To allocate scarce memory resources among competing processes
  - Virtual memory larger than physical memory (Chap. 9)
- ✓ To provide isolation between processes
  - Memory protection

## ■ Virtual memory concept

- ✓ Separation of local address and physical address
- ✓ Translation of logical address to physical address through MMU (Memory Management Unit) in CPU
- ✓ Swapping of memory chunks between main memory and secondary storage
- ✓ Checking of valid references

## ■ Memory management methods

- ✓ Contiguous allocation
- ✓ Paging
- ✓ Segmentation
- ✓ Paging with segmentation

## ■ Contiguous allocation

- ✓ Physical addresses of a process are allocated contiguously
- ✓ Allocation algorithms: first/best/worst-fit
- ✓ Advantages
  - Simple translation
  - Simple memory protection
- ✓ Disadvantages
  - Low memory utilization due to fragmentation ( $\rightarrow$  compaction)



## ■ Paging

- ✓ Motivation
  - Physical addresses of a process can be noncontiguous
- ✓ Address translation
  - Frame: fixed-sized block in physical memory
  - Page: fixed-sized block in logical memory
  - Logical address (p, d) is translated to physical address (f, d)
- ✓ Page table: TLB (Translation Look-aside Buffer)
  - Page to frame mapping table:  $p \rightarrow f$
  - Associative memory (or Content-Addressable Memory)
- ✓ Protection
  - Valid or invalid bit in page table
- ✓ Page table implementation
  - Hierarchical paging
  - Hashed page tables
  - Inverted page tables
- ✓ Easy to share pages



## ■ Segmentation

- ✓ Motivation
  - A program is a collection of segments such as code, data, stack, function, object, ...
  - Segment instead of page in paging
- ✓ Address translation
  - Logical address (s, d) is translated to physical address (b+d)
- ✓ Segment table
  - Segment to base address mapping table: s → b
- ✓ Protection
  - Valid or invalid bit in segment table
  - Check if  $(b+d) < \text{limit}$
- ✓ Easy to share segments

## ■ Paged segmentation

- ✓ Segmentation with paging
- ✓ Paging + Segmentation

